



TECHNISCHE UNIVERSITÄT BERLIN

BACHELOR THESIS

Key Management and Key Exchange for Client-side Encrypted Cloud Storage

Valentin Franck, 364066

Supervised by
Prof. Dr. Stefan TAI

Assistant supervisor
Prof. Dr. Sahin ALBAYRAK

March 16, 2018

Eidesstattliche Versicherung

Hiermit erkläre ich, dass ich die vorliegende Arbeit selbstständig und eigenhändig sowie ohne unerlaubte fremde Hilfe und ausschließlich unter Verwendung der aufgeführten Quellen und Hilfsmittel angefertigt habe.

Berlin, den 16. März 2018.

.....
Valentin Franck

Abstract

Cloud storage systems have become very popular rising questions about the security of the data shared in the cloud. This thesis presents a secure cloud storage system, that uses client-side encryption to protect the confidentiality and integrity of the data in the cloud. It proposes a cryptographic tree structure that enables efficient and secure sharing of the data by using hybrid encryption. The cloud storage provider runs a Public Key Infrastructure, so that other user's public keys can easily be retrieved and used to share data with them. To avoid putting too much trust in the cloud storage provider two certificate verification mechanisms are introduced: QR-Code verification and verification via Socialist Millionaires' Protocol, which can be performed over a potentially insecure channel. In order to add further devices to one's account or recover from a device loss the user's private key is stored in the cloud encrypted with a random generated 12 English words passphrase. The feasibility of the presented design is evidenced by a proof of concept implementation and an evaluation.

Zusammenfassung

Cloudspeichersysteme sind sehr populär geworden, was Fragen nach der Sicherheit der Daten in der Cloud aufgeworfen hat. Diese Abschlussarbeit präsentiert ein Cloudspeichersystem, das Client-seitige Verschlüsselung zum Schutz der Vertraulichkeit und Integrität der Daten in der Cloud einsetzt. Sie schlägt eine kryptographische Baumstruktur vor, die durch den Einsatz hybrider Verschlüsselung, effizientes und sicheres Teilen der Daten in der Cloud ermöglicht. Der Cloudspeicheranbieter betreibt eine Public Key Infrastructure, so dass die öffentlichen Schlüssel anderer Nutzer_innen jederzeit abgerufen werden können, um Daten in der Cloud zu teilen. Um das Vertrauen in den Cloudspeicheranbieter zu minimieren werden zwei Mechanismen zur Zertifikatverifizierung eingeführt: QR-Code Verifizierung und Verifizierung mittels des Socialist Millionaires' Protokoll, das über einen potentiell unsicheren Kanal abgewickelt werden kann. Um ihrem Account weitere Geräte hinzuzufügen bzw. wieder Zugriff auf den Account nach dem Verlust eines Gerätes zu erhalten, werden die privaten Schlüssel von Nutzer_innen mit einer Passphrase aus 12 zufallsgenerierten englischen Wörtern verschlüsselt und in der Cloud gespeichert. Die Umsetzbarkeit des vorgestellten Designs wird durch eine Proof of Concept Implementierung mit Auswertung sichergestellt.

Contents

1	Introduction	1
2	Requirements for a secure cloud storage	4
2.1	Application scenario and desired functionality	4
2.2	Basic security goals	5
2.3	Threat model	5
2.4	Usability	7
3	Cryptographic basics and security infrastructures	8
3.1	Symmetric-key cryptography	8
3.2	Password-based key derivation functions	11
3.3	Public-Key Cryptography	12
3.4	Challenge-Response-Protocols	14
3.5	Public key verification	15
4	Related Work and similar projects	20
4.1	Research on cryptographic cloud storages	20
4.2	Comparison of existing client-side encrypted cloud storage systems . .	22
5	Key Management	30
5.1	Architecture	30
5.2	Operations	35
6	Cross-device Key Exchange	39
6.1	Adding the initial device	40
6.2	Adding further devices	40
6.3	Revoking a X.509 certificate	41
6.4	Sharing data with another user	42
7	Public key verification	44
7.1	Man-in-the-middle attacks	45
7.2	Why the CSP should not be trusted as a CA	45
7.3	Smart Card based Public Key Infrastructures	46
7.4	QR Codes for Key Exchange and Verification	47
7.5	Socialist Millionaires' Protocol	48
7.6	X.509 Certificate and Verification List	49
8	Web Access	49
9	Implementation of the Prototype	51
9.1	Server	51
9.2	Client	52
10	Evaluation of the Prototype	56
11	Conclusion	57

List of Figures

1	An example of a Read Crypttree	32
2	An example of a Write Crypttree	34
3	Sharing data activity diagram	43
4	QR-Code verification activity diagram	54
5	SMP verification sequence diagram	55

1 Introduction

In recent years cloud computing has become a huge hype in computer science. Both companies and private users have started to make use of the cloud in various ways - be it to outsource computing power, hardware maintenance or storage. One famous application provided by cloud computing technologies is reliable data storage.¹ In cloud storage systems users can store their files online, access them from any location and device while at the same time being able to share them with others. This set of functionality is also known as Sync & Share. In spite of the obvious convenience of cloud storage for many use cases (such as collaborative work on a larger set of documents, sharing or syncing (large) files or directories or a simple data backup) questions about privacy and security of the data processed and stored by the cloud storage provider (CSP) have been rising and preventing some users from actually using these services to their full potential.

Famous commercial CSPs, e.g. Dropbox and Google Drive, claim to be secure. Security for these CSPs mainly refers to their use of server-side encryption of the data at rest at the host and the use of transport encryption (SSL/TLS) between client and server and of course between different servers. While these security measures can provide some protection to keep data confidential, eventually it is based on the user's trust in the CSP itself. Since with server-side encryption the CSP (or its employees) may always access both encrypted data and the keys required to decrypt it, the CSP is free to read or even modify stored data at all times. The CSP may also pass on these rights without the user ever acknowledging it. Governments may legally force the CSP to provide customer data and of course in case of a server breach the users' data might also be compromised, since the encryption keys are right there next to the data.

Edward Snowden's revelations, however, have brought some clarity about the actual extent to which the US foreign security agency NSA (among other organizations) applies techniques of mass surveillance to online communication. Snowden has made it very clear that the extensive use of state-of-the-art cryptography is the only way to protect one's privacy online. Since it is not only governments but also private companies (with the intention to analyze the users' data and behavior to improve advertisement) and for example criminal hackers (deploying ransom ware etc.) that might seek access to the data stored in the cloud, server-side encryption is not enough to be counted as an appropriate security measure. Groups of activists, lawyers and journalists, that heavily rely on the confidentiality of their communication are particular targets for state surveillance, which may in some occasions put them at great risk when using insecure standard CSPs. If they are aware of the threat this deprives them from using convenient communication tools and thus is a major hindrance to their work. So particularly targeted groups are in real need of tools that allow them to sync and share their data and resolve the mentioned shortcomings regarding security of the most popular CSPs.²

¹E.g. even the Amazon S3 Reduced Redundancy Storage stores data various times redundantly on different machines and thus achieves a much higher availability than a traditional backup on a hard disk (Amazon 2018).

²Mobile Messengers such as WhatsApp have reacted to that threat by applying the Signal Protocol, a modern cryptographic protocol for asynchronous messaging, that provides "end-to-end

Current client-side encrypted³ cloud storages suffer from several issues. In some cases encryption is realized using containers instead of using different keys for each file. That means, the whole container has to be re-encrypted and uploaded again, even if the user just modified one single file. More importantly there are severe security problems. Some solutions do not allow users to share data at all, while others that do, require users to trust the CSP, that data is really only just shared with the person the CSP claims. The problem arises, because these CSPs use a Public Key Infrastructure (PKI), in which they act both as Certificate Authorities (CA) and key distributors. So in this scenario, if Alice wants to share a file with Bob, she has to trust the CSP, that she really received Bob’s public key and that the CSP is not performing a man in the middle attack. But there is another problem. Either users are not able to use multiple devices, because they cannot transfer their private keys or the server is responsible for the transfer. Since these private keys are at best encrypted with the password of the user account (which is usually relatively weak), the server might even get direct access to the user’s private key and thus to all data she stores in the cloud.

This thesis presents solutions to all three problems. We use a cryptographic tree structure as first proposed by Grolmund et al. (2006), which resembles the folder structure. In this tree structure every file and every folder is encrypted with a new key. Every file and folder key is encrypted with the parent directories folder key up to the user’s root directory. The key for the root directory is encrypted with the users private key. In order to share a folder with Bob Alice simply has to decrypt the folder key and re-encrypt it with Bob’s public key. Bob can then use the folder key to recursively decrypt all contained file and subfolder keys. This protects the confidentiality of the data. A second cryptographic tree structure is introduced to distinguish between read and write access. In this second write tree a private key is stored for every file and every folder. Whenever Alice or Bob modify a file, they use this key to sign the new version, so that anybody with read access can verify the signature, because the corresponding public key is stored in the read tree. By using digital signatures we protect the authenticity and integrity of the shared data.

The solution we present in this thesis does not require any trust in the CSP. As mentioned before the CSP provides a PKI, in which it is the CA and we also make leverage of the availability of the CSP for certificate distribution. So whenever Alice wants to share a file with Bob she can retrieve Bob’s certificate issued by the CSP. Alice is now technically able to share data with Bob. Too make sure Alice really has Bob’s valid certificate and not one controlled by the CSP, we provide

security, deniability, forward secrecy, and future secrecy,” (Frosch et al. 2016) while not making concessions to user experience. The Signal Protocol has first been developed by OpenWhisper Systems for TextSecure, which is nowadays simply called Signal.

³In this work we will use the terms client-side and end-to-end encryption mostly synonymously in the context of cloud storage. By both we refer to a system that encrypts the users’ data on the users’ devices before uploading it to the server. The server does not get access to the keys in plaintext and therefore cannot access the data either. Sometimes client-side encryption is referred to as a system that applies both end-to-end encryption and local encryption of the data at rest at the client. The term end-to-end encryption originally comes from software more focused on the exchange of messages such as email and instant messaging. Since there are plenty of solutions for local encryption, that can easily be implemented on top of our system, we will not go into detail about local encryption.

two possible verification mechanisms: QR-Code verification and verification via the Socialist Millionaires' Protocol (SMP). For QR-Code verification Alice and Bob have to meet or Bob has to send Alice the QR-Code generated on his device. Alice can scan the QR-Code and if it succeeded, Alice and Bob have mutually verified that they are really both using each other's certificates. The SMP on the other hand is a cryptographic protocol that allows Alice and Bob to mutually verify their certificates using a potentially insecure in-band channel. They only need to share a weak natural language secret, like the name of the restaurant, they had dinner at on their last vacation.

In order to allow users to recover from private key loss and give new devices access to their account, users have to transfer the private key between devices. For this purpose we randomly select 12 words from a 2048 English words list. We then apply a key derivation function and use the resulting key to symmetrically encrypt the private key. In this encrypted form the private key is stored in the cloud. So if Alice wants to add a new device she just downloads her own private key from the cloud storage and uses the passphrase only she knows to decrypt it. She then has gained full access to her account. Using this 12 word passphrase enforces passwords with enough entropy while not requiring complicated interaction from the user.

With our application scenario in mind it is important that we achieve our security goals without sacrificing usability, i.e. the Sync and Share functionality should not be preserved. Our proof of concept implementation shows that we achieved this goal with only one notable yet deliberate exception. Web access is not supported by the design we chose. Since there are currently no easy ways to provide web access in a truly secure manner, we decided not to include it, but it could easily be added for convenience sacrificing some of the achieved degree of security, as we will explain in section 8. Apart from usability we consider it a crucial property of our software to be open-source (or at least have frequent external security audits), because otherwise users will never be able to confirm that the encryption actually works as stated by the CSP and that there are no back doors, which might be used to access the confidential data store by high profile users such as activists.⁴

In the following sections we will further elaborate our system, beginning in section 2 requirements for a cloud storage system and especially our threat model. In section 3 we will explain the cryptographic basics, because the security heavily depends on using secure algorithms and implementing them correctly in software. Section 4 analyzes related work and also provides an overview over other client-side encrypted cloud storages analyzing their strengths and weaknesses. In the next three sections we present the design choices for our own design as drafted above. Section 5 explains the cryptographic tree structures that enable flexible sharing, modification of files and revocation of shares. The details of how key exchange works are described in section 6. In section 7 we discuss possible verification mechanisms and present the ones we finally decided are best suitable. We finish the system design part in section 8 explaining why web access and sharing via encrypted web links is not supported by the proposed system, because it would be a security threat. Finally in section

⁴Note that the open-source cloud storage software Nextcloud has introduced end-to-end encryption during the elaboration of this thesis. It is a good approach, but unfortunately does not provide optimal usability and lacks proper key verification mechanisms (see section 4.2.6).

9 we present the details of our proof of concept implementation and evaluate it in section 10.

2 Requirements for a secure cloud storage

In this section we will briefly explain the application scenario and use cases the cloud storage system proposed by this work is designed for. We will then develop a threat model and clarify the core security goals of our design. We will finish this section with a short explanation, why usability is a crucial property for our design (and should be for any secure privacy preserving software).

2.1 Application scenario and desired functionality

The application scenario that we have in mind is to provide secure cloud storage to private users and small-scale organizations, such as Non-Governmental-Organizations (NGOs), small groups of activists, journalists, lawyers and other associations, that are particular targets of (state) surveillance and repression. This means the focus of the design is rather on privacy and security than performance and scalability, as long as the system remains usable.

Typical use cases for the system are as follows:

- **Data storage and backup:** The user uploads personal data, in order to keep a backup or to be able to access it via a web panel from her web browser. This can be anything from a single documents to a complete hard drive image.
- **Syncing data:** The user can use the cloud storage system to automatically sync data between several of her devices.
- **Sharing data:** The user uploads data, in order to share it with other users. This may be single files or complete hierarchical directories. The user wants it to be as easy as possible to decide which other users get access to the data. The user determines, what kind of access rights other users are granted. Currently the system supports to differentiate between read and write access but this could be further differentiated.⁵ The user also wants to be able to revoke these access rights at all times.

These are just the most common uses cases for the intended users of the system and this list does not pretend to be exhaustive. Nevertheless this it is enough to give us an idea about how cloud storages are used and what the consequences of a compromise of confidential data might be.

The system might be run on a private or on a public cloud, e.g. small groups of activists might even use a shared host to run the cloud storage server, while others might rely on public cloud services such as Amazon S3, because it provides elastic and reliable storage. Note however, that the latter mostly affects the server-side of the system, which is only part of this work, whenever it is needed to ensure the usability of client-side encryption.

⁵These access rights are cryptographically secured by our design, while an ACL on the server might easily allow for more options.

2.2 Basic security goals

There are several core concepts in information security referring to different protection goals a secure system should comply. The best known is the “CIA triad”, confidentiality, integrity and availability. For our context we will add authenticity and reliability as key properties.

- *Confidentiality*: No one, except a legitimate recipient, is able to learn any information about a user’s data.
- *Integrity*: Any unauthorized modification of the user’s data (including creation and deletion) will be detected (and at best can be rolled back).
- *Availability*: The user’s data is (*efficiently*) accessible at all times from any machine by an authorized user.
- *Authenticity*: Any communicating entity really is, who it claims to be and the author of any modification of the system can be identified.
- *Reliability*: The system works as intended and with consistent results.

In this thesis we will not answer the question, how to protect reliability and availability due to the chosen threat model, in which we assume that the server will not attack availability of the users’ data.

Because it is crucial “that the security properties described above are achieved based on strong cryptographic guarantees as opposed to legal, physical and access control mechanisms” (Kamara and Lauter 2010, 137), in the following section we will take a closer look at the required cryptographic basics. Kamara and Lauter (2010, 138) also point out that “it is important that the client-side application and, in particular, the core components be either open-source or implemented or verified by someone other than the cloud service provider.” This is why we also make it a key requirement for the client to be open-source.

2.3 Threat model

In this section we will develop a threat model. Credit is given to Mathieu Bourrier (2015), on whose threat model for a proposed extension to the ownCloud Desktop client to add client-side encryption, we built our work. First of all we define that any of the user’s or a friend’s devices (and the client it runs) is trusted. The user needs to trust her own devices and also those of her friends in order to share data with them. If the user’s device was compromised, it could easily, sent all plaintext to an attacker (cf. Wilson and Ateniese 2014, 406). That means, if the user’s account is compromised or data is compromised on a trusted device, due to local vulnerabilities, this is not considered a break of the security model and any user concerned about the security of her data should take further security measures, such as local file encryption, system hardening etc.⁶ However, a user’s account should not be compromised due to an insecure cloud storage system.

⁶The current debate on source surveillance in telecommunication has made it clear, that with the increasing use of transport encryption and end-to-end encryption, governments are willing to

Networks, the data passes on its way between server and client are considered insecure and attackers will actively try to listen in or even modify data on these networks. This problem can easily be resolved with transport encryption (SSL/TLS) (cf. Bourrier 2015, 5) and is therefore not part of our work.

In a typical scenario we assume that an attacker wants to get read access to the users' data stored on the cloud storage system. The attacker, the proposed system pretends to protect against, is a government institution trying to read the documents stored and shared by activists, journalists and lawyers on the cloud storage system. This means the attacker has many (legal) resources. We assume, that the attacker will mostly try to remain covert, as governments usually do, when monitoring their surveillance targets. Note, that this is not necessarily the case and in some situations attackers might also try to manipulate the stored data or in case of subpoenas, the attacker might openly try to get access to the stored data. That means the CSP cannot be trusted to protect and respect the confidentiality of the users' data. It will, however, stick to the protocol and most importantly protect the availability of the users' data. Furthermore for the case of a server breach, an attacker with access to the host might try to modify user data, if she cannot access the plaintext. This means a government will mostly act as an honest but curious attacker, that might have access to the CSP.

While it has become a standard not to trust the network, not trusting the host requires security measures far more elaborated than transport encryption with SSL/TLS. If an attacker controls the host, she should not be able to access or modify any user's plaintext data without knowledge of the user. But the attacker is allowed to be able to get some information on the metadata, e.g. an estimation of how much data there is and of how many files it consists of and who accesses the data.⁷ The attacker could also delete or prevent access to the data, but this is not assumed to be her intention. Therefore availability of the data is not a primary security objective in our scenario and the CSP is trusted to take care of it.⁸ An attacker controlling the host might also be able to provide outdated versions of files to users. This might even include providing different versions to different users (Cachin, Keidar, and Shraer 2009). However, this is an unlikely scenario with relatively low risks for the assets since it requires a lot of effort from the attacker and the system therefore does not protect against such an attack. It is much more likely that infrastructures are completely shutdown as done by authoritarian governments in the past.

The key asset is the user data, which under all circumstances must remain confidential. An unauthorized access to the stored data is considered to have the most severe consequences and therefore represents the greatest risk for an activist. Our

gain direct access to the communication data on the source devices and they are currently working on the legal framework, that gives them a lot of competences to infiltrate devices for surveillance (Meister, Andre 2017).

⁷Some of these things, such as analyzing user cooperation might possibly be prevented using modern functional cryptography. However they are not the focus of our scenario. For further reading see: Kamara and Lauter 2010.

⁸Availability and reliability are some of the main reasons users migrate their data to cloud storages. It is therefore highly unlikely that a cloud storage provider does not protect availability. Our main focus is on the kind of security violations that happen unnoticed and can therefore be particularly harmful.

system protects the integrity, confidentiality and authenticity of the data stored in the cloud and in transmission. It does not protect its availability on the cloud and does not protect the data stored on the users' devices at all.

2.4 Usability

In their case study Whitten and Tygar (1998) evaluate the usability of security software and come to a clear conclusion: The more efforts a technology requires users to make, the less likely they are going to adopt it. In software engineering usability “is a quality attribute that assesses how easy user interfaces are to use. The word ‘usability’ also refers to methods for improving ease-of-use during the design process. Usability is defined by 5 quality components” (Nielsen 2017):

- *“Learnability:* How easy is it for users to accomplish basic tasks the first time they encounter the design?
- *Efficiency:* Once users have learned the design, how quickly can they perform tasks?
- *Memorability:* When users return to the design after a period of not using it, how easily can they reestablish proficiency?
- *Errors:* How many errors do users make, how severe are these errors, and how easily can they recover from the errors?
- *Satisfaction:* How pleasant is it to use the design?”

While utility describes, whether a software provides needed functionality, usability describes how easy it is to actually benefit from these features (Nielsen 2017). The utility of a client-side encrypted cloud storage system has been sufficiently described. Therefore in order to make it useful we need to make sure that the proposed system provides good usability. In a market with plenty of CSPs available, solutions that do not provide good usability will be discarded. To achieve our goals and promote the general use of privacy preserving cloud storage systems usability is a key factor. Usability, is sometimes used interchangeably with the term user experience, which denotes the emotional and more subjective experience a user makes while interacting with the software. Usability is usually considered a prerequisite for a good user experience, because no user will ever make a positive experience, if she cannot get her task done (efficiently) (Law et al. 2009).

In this thesis we do not want to go into the details of the research that has been done in the field of usability, but rather take a practical approach to it. We will consider, that the existing cloud storage solutions (to a slightly varying degree) already provide good user experience. So our main goal will be to conserve the usability of these systems. We will therefore make the encryption fully transparent where possible and make usability a priority in any design decision. However in certain circumstances full transparency might not be possible to achieve without sacrificing the gain in security our software aims for. In these cases we will limit the additional user interaction to a minimum and make it as easy to understand as

possible.⁹ By transparency we mean, that the user will not be able to tell whether the files are actually encrypted, because no interaction from her part is required and she can use the cloud storage, just as if there was no encryption at all.

For the most part the system we propose in this paper will therefore not sacrifice user experience. However some features might be hard to implement in a fully client-side encrypted cloud storage system. Examples are web access and the option to recover one's account after a password (or key) loss.

3 Cryptographic basics and security infrastructures

In this section we will briefly discuss the cryptographic basics required to design a secure cloud storage system. We will sketch out how the discussed protocols work, what their pitfalls are and whether they are currently considered secure. The focus will be on those protocols and algorithms, that will later be suggested for use in a client-side encrypted cloud storage system.

3.1 Symmetric-key cryptography

Symmetric-key Cryptography is the oldest form of cryptography. It is still widely used because of its relative simplicity and far better performance than its counterpart public-key cryptography. It is also widely used in hybrid cryptographic systems, in order to benefit from both the performance of symmetric-key cryptography and the easement of the key negotiation or exchange provided by public-key cryptography. Modern symmetric-key cryptography is not only used for encryption/decryption, but also used in cryptographic hash functions. In symmetric-key encryption the communicating parties, say Alice and Bob, first have to agree on a shared key, which they keep secret. Alice uses that key to encrypt a message to Bob, who uses the same key to decrypt it. Although the algorithms for encryption and decryption are publicly known nobody should be able to get information about the plaintext of a message from its ciphertext without knowledge of the used key. Because both parties use the same key for encryption and decryption we speak of symmetric-key encryption. The main challenge in this scheme, how to exchange the shared key, has led to the invention of public-key cryptography.

“Among all encryption algorithms, symmetric-key encryption algorithms have the fastest implementations in hardware and software. Therefore, they are very well-suited to the encryption of large amounts of data” (Delfs and Knebl 2015, 11f).

There are basically two different types of symmetric encryption algorithms:

“Block ciphers operate on data with a fixed transformation on large blocks of plaintext data; stream ciphers operate with a time-varying transformation on individual plaintext digits” (Rueppel 1992).

⁹One use case, in which full transparency will be hard to achieve, is public key verification.

In this paper we will only explain one block cipher algorithm, AES, since block ciphers are most widely used in software today, because they can be implemented more efficiently in software (Schneier 1996, 211).

3.1.1 Advanced Encryption Standard (AES)

AES is a relatively new encryption algorithm, yet widely regarded as exemplary in its creation. Since its predecessor DES was not regarded to be entirely secure anymore the Nation Institute of Standards and Technology (NIST) started a selection process in 1997 for the new Advanced Encryption Standard (AES). Proposals could be made and finally three years later the Rijndael algorithm was chosen in a public process. AES is an iterated block cipher, that has a fixed block length of 128 bits and key length of 128, 192 or 256 bits. Depending on the key length the algorithm applies a round function 10, 12 or 14 times after one initial round key addition. The algorithm first creates a state, represented as a matrix and then round wisely applies byte substitution, row shifting and column mixing. However the only operation depending on the secret key is the addition (XORing) of the initial key and of a round key in every round. Those keys are derived from the secret key by a key expansion function.

The algorithm is obviously invertible, because XORing can be inverted by simply applying it twice and the other operations can simply be reverted. The round function, however, ensures some crucial properties of AES. The byte substitution causes nonlinearity (confusion between ciphertext and secret key) and the row shifting provides diffusion. i.e. every bit change in the key results in a 50% chance of a bit change in every bit of the ciphertext and every bit change in the plaintext results in a 50% chance of a bit change in every bit of the ciphertext (Ferguson 2010, 55). Although the focus of AES was mainly on performance and not security, it is still reasonably secure to use. However, there have been certain advances towards breaking AES:

“The attacks against the full 192- and 256-bit versions of AES are theoretical—not practical—so we aren’t ready to lose any sleep over them just yet. [...] Given all that we know today, using AES still seems like reasonable decision. It is the U.S. government standard, which means a great deal. Using the standard avoids a number of discussions and problems. But it is important to realize it is still possible that future cryptanalytic advances may uncover even more serious weaknesses” (Ferguson 2010, 55f).

However, until today none these first approaches to break AES has led to success in practice which is why is why Pancholi and Patel (2016) consider that the “AES algorithm is a highly secure encryption method”, that has a “high performance without any weaknesses and limitations.”

Due to its great performance and security we will use AES for symmetric encryption in the secure cloud storage proposed by this thesis. The main advantage of AES is that it is “the standard” and therefore supported by most cryptography libraries (Ferguson 2010, 59).

3.1.2 Hash functions

Cryptographic hash functions are one application of symmetric cryptography. A hash function is a one-way function, which means that it is easy to compute the hash value $H(M)$ of any given message M , while it is practically not possible to compute the original message M just by knowing its hash value $H(M)$ and the hash function H . Given M it should also be hard to find another message M' with $M \neq M'$ and $H(M) = H(M')$. For some applications collision-resistance is also needed: it is practically not possible to find two messages $M \neq M'$ with $H(M) = H(M')$ (Schneier 1996, 429).¹⁰

3.1.3 Keyed-Hash Message Authentication Code (HMAC)

When cryptographic hash functions are used for message authentication, they are called Message Authentication Codes (MAC). In this case a hash function is used to compute a unique fingerprint (of fixed length) of a given message. This hash value is encrypted with a secret key.¹¹ Since it is impossible to change the message and recompute a valid MAC without knowledge of the key, MACs ensure the integrity and authenticity of the message. A MAC can also only be verified by someone with access to the secret key (Schneier 1996, 455f).

The most commonly used method today to create a MAC from a hash function is the keyed-hash MAC (HMAC). HMAC does not specify one particular hash function, but can be used with different hash functions such as MD5 or the SHA-family. To guarantee the security of the MAC, the hash function is applied twice.¹² The HMAC of a message M is computed as follows:

$HMAC(K, M) = H((K \oplus opad) || H((K \oplus ipad) || M))$ where K is the key, H is the used hash function and $ipad$ and $opad$ are the constant inner ($0x36$ repeatedly) and outer padding ($0x5C$ repeatedly) used to XOR the key with. $ipad$ and $opad$ are used to derive two different keys for the concatenation with the message and its hash value respectively in the inner and outer application of the hash function (Delfs and Knebl 2015, 42ff).

One of the advantages of HMAC over traditional MACs is the variability of the length of the resulting HMAC depending on the used hash function. It is easy for

¹⁰This is a stronger requirement than the one mentioned before and needed to avoid birthday attacks (Schneier 1996, 165f). Some cryptographic hash functions such as MD5 are no longer considered to be collision-resistant and should therefore not be used in digital signatures and other critical applications. In 2005 the Chinese scientists Wang, Yin and Yu found a way to ease the finding of collisions in SHA-1. However, the improved algorithms of SHA-2 (with 256, 384 and 512 bit hash value length) are still considered to be collision-resistant. Encouraged by the great success of AES the NIST started a similar selection process for a new secure cryptographic hash algorithm SHA-3 in 2007. In 2012 Keccak's algorithm was chosen. Because SHA-2 is still viewed as secure and is included in more cryptographic libraries, the use of SHA-2 is suggested in this thesis (Delfs and Knebl 2015, 40ff).

¹¹If the hash value is not encrypted with a shared secret key, but with a private asymmetric key, it is a digital signature. Everyone with knowledge of the corresponding public key can then verify, that the signature has been made by the owner of the private key.

¹²This prevents the length extension attack, where an attacker with knowledge of $MAC(M)$ is also able to compute the MAC for another message M' attached to M : $MAC(M || M')$ (Swoboda, Spitz, and Pramateftakis 2008, 107).

the receiver of the HMAC to verify, that the sender knows the shared secret key and the message has not been modified. On the other hand it not possible to find a message M for a given HMAC, due to the one-way property of cryptographic hash functions. This is indeed possible for other MACs under certain circumstances (Swoboda, Spitz, and Pramateftakis 2008, 108). Due to these characteristics HMAC can be viewed as a secure way to ensure authenticity and integrity in messages provided that the communicating parties share a secret key.

3.2 Password-based key derivation functions

In many security applications it is not possible (or at least not desirable regarding usability) that the user always provides her secret key. While long secret keys are impossible to remember or to transfer to another device, passwords are not. Password-based key derivation functions have therefore been developed to derive secret keys from a user's password. One use case of password-based key derivation functions is calculating a hash value of the password that is hard to attack using brute force, rainbow tables or dictionary attacks.¹³ In this section we will briefly discuss, how PBKDF2 works and what degree of security it can provide. The problem with passwords being the basis of a secret key is obvious:

“[P]asswords, in particular those chosen by a user, often are short or have low entropy. Therefore special treatment is required in the key derivation process to defend against exhaustive key search attacks” (Yao and Yin 2005).

3.2.1 PBKDF

Generally key derivation functions use the following scheme to derive the key K : $K = H^C(P||S)$. H is a function, that is applied C times to the concatenation of the password P and the salt S . In the industry standard PKCS#5 two different functions PBKDF1 and PBKDF2 are defined. H in PBKDF1 is a hash function such as MD5 or SHA-1 and the salt S should be at least 64 bits long, while the number of iterations C should at least be 1000. PBKDF2 is designed to be more secure and uses a keyed-hash function, such as HMAC, instead of a simple hash function. The password in this case is used as the key. Adding some extra protection by XORing the results of each iteration the key is derived as follows: $K = H_P^1(s) \oplus H_P^2(s) \oplus \dots \oplus H_P^C(s)$ A high number of iterations makes sure that it is costly to compute the key from a password. By the concatenation with a salt the entropy of the password is further increased, which makes it practically infeasible to build rainbow tables. The application of both methods therefore significantly decreases the success chances of brute force and dictionary attacks. In their paper Yao and Yin (2005) show that the mechanisms used in PBKDF1 and PBKDF2 indeed make these functions secure in case that the attacker cannot choose the parameters such as the salt and the number of iterations. Otherwise these functions have serious vulnerabilities. Colin Percival, the inventor of the password-based key derivation function Scrypt, makes a strong case not to

¹³These are the only three feasible attacks against modern password-based key derivation functions, because they are based on strong hash functions (Percival 2009, 1).

put too much trust into a function like PBKDF2, since users tend to choose weak (low-entropy) passwords:

“Unfortunately, this form of “brute force” attack is quite liable to succeed. Users often select passwords which have far less entropy than is typically required of cryptographic keys; a recent study found that even for web sites such as paypal.com, where — since accounts are often linked to credit cards and bank accounts — one would expect users to make an effort to use strong passwords, the average password has an estimated entropy of 42.02 bits, while only a very small fraction had more than 64 bits of entropy [15].” (Percival 2009, 1)

The use of password-based key derivation functions should therefore be considered carefully for each particular application scenario, because otherwise users are able to sacrifice the security of the complete system.

3.3 Public-Key Cryptography

The invention of public-key cryptography was one of the most important advances in cryptographic methods that was ever made. It was first introduced in the mid 1970’s by Diffie and Hellman (1976), who published their still widely used key exchange algorithm and then by Rivest, Shamir, and Adleman (1978), who published the algorithm today best known by their initials: RSA. Prior to the invention of public-key cryptography it was impossible to exchange or negotiate symmetric keys over an insecure channel. Also, each pair of communicating parties needed their own secret key, which meant that it did not scale for large communicating groups.¹⁴

The basic idea of public-key cryptography is that each communicating entity owns two keys. One that is public and accessible for everyone. The other one that is private and must always remain the secret of this entity. If Alice wants to send a message to Bob, she uses Bob’s public key to encrypt it and then sends it to him. On receiving Alice’s message Bob uses his private key to decrypt it. That means, that anyone can encrypt a message using Bob’s public key, but only someone who possesses the corresponding private key, that is Bob, will be able to decrypt it. Of course it is crucial, that knowledge of the ciphertext, the plaintext of the message and the public key used for encryption do not allow any conclusions about the private key, since this would compromise all other messages encrypted with this key pair. It must also be infeasible to get any useful information on the plaintext of a message by mere knowledge of its ciphertext and the public encryption key. The first algorithm, that allowed such characteristics was RSA. Later, however, certain improvements were made by similar approaches such as elliptic-curve cryptography, which is based on a similar mathematical problem as RSA.

Since public-key cryptography is far slower and uses a lot longer key lengths than conventional symmetric-key cryptography, it was not immediately clear, how crypto systems could benefit from public-key cryptography:

“In viewing public-key cryptography as a new form of crypto system rather than a new form of key management, we set the stage for criticism

¹⁴If there are n communicating parties, the number of keys to be exchanged is in $O(n^2)$.

on grounds of both security and performance. Opponents were quick to point out that the RSA system ran about one-thousandth as fast as DES and required keys about ten times as large. Although it had been obvious from the beginning that the use of public key systems could be limited to exchanging keys for conventional cryptography, it was not immediately clear that this was necessary. In this context, the proposal to build *hybrid* systems was hailed as a discovery in its own right” (Diffie 1992).

This means that while symmetric-key cryptography is best suited to encrypt data due to its performance and security, public-key cryptography should be used for key management and key exchange in particular, thus making a hybrid system the most secure and effective we can think of today. In a hybrid system the data is encrypted with a random symmetric key, that in turn is encrypted with the receiver’s public key. Today public-key cryptography has its well-established field of application, that does exceed encryption, e.g. it is also used for digital signatures, public-key certificates etc. (cf. Schneier 1996, 216, 461).

3.3.1 RSA

Of all public-key algorithms RSA is probably the most popular. This is on the one hand due to its early publication and on the other hand due to the ease of understanding and implementation it provides (Schneier 1996, 466f). It is still used in a wide number of applications such as PGP or the Diffie-Hellman Key Exchange in SSL/TLS. However, nowadays elliptic curves have replaced RSA in many applications because of their efficiency and the relatively small length of the required public keys.

The security of RSA is based on the difficulty to factorize large numbers. For key generation the user chooses two large prime numbers p and q of equal length (in bit representation) and then computes the product $n = pq$. The (public) encryption key is randomly chosen “such that e and $(p-1)(q-1)$ are relatively prime” (Schneier 1996, 467). The decryption key d is computed with the Euclidean Algorithm: $ed \equiv 1 \bmod (p-1)(q-1)$. This means that d and n are also relatively prime. The public key thus consists of the tuple (e, n) and the private key of the tuple (d, n) . For encryption a message m is divided into blocks of a fixed length, depending on the length of p and q . The ciphertext of the i -th block is computed as follows: $E(m_i) = c_i = m_i^e \bmod n$. Decryption is just the inversion of the encryption operation: $D(c_i) = m_i = c_i^d \bmod n$ (Schneier 1996, 466ff).¹⁵

RSA, if implemented correctly, is still a secure algorithm, since the best attack mathematically known is to factorize n . However it has not been *proven* that the security of RSA is based on the factorization of large numbers. So, theoretically, there might be an easier attack. The factorization of large numbers, however, is believed to be a hard problem, that has no polynomial solution. This has not been proved either (Ertel 2012, 82f). A chosen plaintext attack on public-key crypto systems is easy. Given a ciphertext c encrypted with Alice’s public key e_A Eve can encrypt plaintext messages m using e_A until she finds one with $E_{e_A}(m) = c$.

¹⁵For a proof that $D(E(m_i)) = E(D(m_i)) = m_i$ see Ertel 2012, 81.

This m is the original message. However this attack can easily be prevented by a good implementation that uses randomized padding before encrypting the message, so that the same plaintext is mapped to different ciphertexts. A chosen ciphertext attack combined with social engineering is also possible, but practically extremely unlikely to be successful, which is why RSA can still be considered a secure algorithm (Ertel 2012, 87).

3.3.2 Digital Signatures

Public key pairs, such as RSA pairs, cannot only be used for encryption but also for digital signatures. If Alice wants to digitally sign a message to Bob, she computes the hash value of the message, encrypts the hash value with her private key and then at best encrypts the message and the signed hash value with Bob’s public key. On receiving Alice’s message Bob decrypts it first with his private key and then decrypts the hash value with Alice’s public key. Bob computes the hash value of the message he received and compares it to the one Alice signed. If they match he has successfully verified the authenticity and integrity of the message he received. This works because of the equation $D(E(m_i)) = E(D(m_i)) = m_i$. Both, public and private key can be used for encryption and decryption, that means that encryption and decryption are inverse functions. However, different key pairs should be used for signatures and encryption (Schneier 1996, 483ff).

3.4 Challenge-Response-Protocols

Zero-Knowledge is a term often used by the marketing of cloud storage providers. SpiderOak and Tresorit have used the term to advertise the use of end-to-end encryption and their privacy policy. But the term it is not always used adequately, since it has a well-defined meaning for cryptographers (SpiderOak 2017c). In this section we will first explain how challenge-response protocols work and then go on to describe the idea of a Zero-Knowledge-Proof. This will also make clear, why the marketing departments of the before mentioned companies used the term in the first place—although not perfectly adequately.

Challenge-Response-Protocols have the purpose to identify one party to another, e.g. if Alice wants to log into Bob’s server, she has to prove her identity to Bob. Instead of simply transmitting a username and password or secret key, a challenge-response protocol is somewhat more elaborated, in order to preserve Alice’s secret, i.e. a password or a key are never transmitted in plaintext to the server. The generic algorithm works like this:

Bob sends a challenge to Alice, e.g. a nonce r . Alice uses that nonce and a function f_k only she and Bob know to compute the response $R = f_k(r)$. She sends Bob her response R . Bob can now verify, that $r = f_k^{-1}(R)$. In this case Alice has successfully authenticated herself to Bob. Alice herself, however, cannot be certain about Bob’s identity.¹⁶ In order to make this protocol work Alice and Bob must have exchanged some information beforehand. This could be a symmetric key and the function could

¹⁶The authentication can be made mutual by a small extension of the protocol (Swoboda, Spitz, and Pramateftakis 2008, 152ff).

be $f_k(r) = \text{HMAC}(r)$. They may also use digital signatures, since only Alice, as the owner¹⁷ of her private key can sign the nonce. However, for Bob it is possible to verify Alice's signature with mere knowledge of her public key. Usually the function will depend on Alice's password (which can be used to derive a key). One problem with this protocol is, that Bob also needs to have access to Alice's secret. That means he has to store her password or her key in his database (Ertel 2012, 102f). In order to prevent this Zero-Knowledge-Protocols have been invented.

3.4.1 Zero-Knowledge-Protocols

Zero-Knowledge-Protocol, or Zero-Knowledge-Proofs as they are also called, are a special type of Challenge-Response-Protocols. They serve the same purpose: Alice identifies herself in her communication with Bob. Alice uses her secret to do that—differently to conventional Challenge-Response-Protocols Bob this time gets no knowledge of Alice's secret. In this scenario Alice is called the prover and Bob the verifier.¹⁸ One example for a Zero-Knowledge-Protocol is the Socialist Millionaires' Protocol, which is used for key verification in Off-the-Record Messaging (see 3.5.4).

A Zero-Knowledge-Proof in the context of authentication has to fulfill the following three properties (Lâm 2016; Delfs and Knebl 2015, 118ff):

- Completeness: If the prover has the claimed identity, the verifier will be convinced by him.
- Soundness: If the prover does not have the claimed identity, he will not convince the verifier (except with some marginal probability).
- Zero-Knowledgeness: The verifier never gets any knowledge of the prover's secret except for the fact that he has the claimed identity

The first two are properties of any Challenge-Response-Protocol, but the third one is new (Schneier 1996, 101ff). There exist possible attacks against Zero-Knowledge-Protocols. However if one secret is only ever used with one communication partner, most of them are rendered useless.¹⁹

3.5 Public key verification

The invention of public-key cryptography has certainly been a crucial step towards today's wide use of cryptography in all kinds of applications. It has made key exchange a lot easier. However, it has brought its own kind of difficulties, as we have seen in the previous section on man-in-the-middle attacks. The challenge for

¹⁷In this thesis we will refer to "ownership" not as the legal concept but rather as having access to a key, data etc. E.g. in the context of cloud storage we will call the user, who initially uploaded data the "owner" of the data.

¹⁸A Challenge-Response-Protocol as described above that uses digital signatures is very close to a true Zero-Knowledge-Protocol. But since Bob as the verifier can collect messages and their signatures he might be able to accumulate some (very marginal) information on Alice's private key, which makes this protocol not Zero-Knowledge (Ertel 2012, 159).

¹⁹For further reading see Schneier 1996, 109ff.

Alice, when sending a message to Bob, is how to get Bob's public key and how to verify that it really is Bob's key and has not been replaced by some attacker, who is pretending to be Bob. Fortunately this problem has been addressed by researchers and several solutions have been proposed.

There are three common ways for Alice to get Bob's public key (cf. Schneier 1996, 185f):

- get it from Bob directly
- get it from a centralized database
- get it from her own database

There are different ways for Alice to verify, that the public key she received really belongs to Bob. The first one is to meet with Bob in person (or call him) and verify that the fingerprint of the key matches. A fingerprint is a hash value, that is unique for every key. The important part about this method is, that an out-of-band channel is used, which is unlikely to be controlled by an attacker. The use of an out-of-band channel is not always possible or practical in many cases. Therefore Webs-of-Trust and Public-Key-Infrastructures have been developed. They are both based on the same idea.

3.5.1 Web of Trust

A Web of Trust works as follows. If Alice cannot make sure personally, that she has Bob's public key, she can still have enough confidence in having the right key, if someone she trusts, tells her it is the right key. In a Web of Trust users sign other users' public keys, once they have verified the keys. So if Alice has verified Bob's public key and Bob has signed Carol's key. Alice can decide that, she will trust Carol's key too, since she trusts Bob and has verified his key. When a lot of users are involved in this and sign each other's keys, this creates a Web of Trust.²⁰ There is one obvious problem with Webs-of-Trust: they require a lot of active contribution and a lot of awareness for security issues from the users. Even if Alice believes Bob would not maliciously sign Carol's wrong key, can she be sure, that Bob put in all the necessary work to verify it was really Carol's key, he signed? Alice certainly can never be sure and will be so even less, if she does not know Bob personally. The second big problem is, that Bob in this scenario actively has to decide, whether he wants to sign Carol's key or not. After signing it he has to upload the signed key to a key distribution server or send it back to Carol. This makes for a terrible user experience (cf. Ertel 2012, 120f).

Both these problems made Webs of Trust an acceptable solution, when they were first introduced at the invention of PGP in the 1990's, but there certainly are better solutions today. Webs-of-Trust cannot provide enough security and are a major obstacle to a good usability for any application.

²⁰Alice of course might also decide not to trust Carol's key, if Bob is the only person, she trusts, that has signed it, i.e. she might require at least two (or more) trustworthy signatures before she trusts another key.

3.5.2 Public key infrastructures (PKI)

Today the most common solution for the problem of key verification are public key infrastructures (PKI). The purpose of a PKI is to allow us to tell, who a public key belongs to.

“There is a central authority that is called the *Certificate Authority*, or CA for short. The CA has public/private key pair (e.g., an RSA key pair) and publishes the public key. We will assume that everybody knows the CA’s public key. [...] Alice generates her own public/private key pair. She keeps the private key secret, and takes the public key PK_A to the CA and says: ‘Hi, I’m Alice and PK_A is my public key.’ The CA verifies that Alice is who she says she is and then signs a digital statement that states something like ‘Key PK_A belongs to Alice.’ This signed statement is called the *certificate*. It certifies that the key belongs to Alice.

If Alice now wants to communicate with Bob, she can send him her public key certificate. Bob has the CA’s public key, so he can verify the signature on the certificate. As long as Bob trusts the CA, he also trusts that PK_A actually belongs to Alice” (Ferguson 2010, 275).

There is not one single CA, but rather several co-existing hierarchies of thousands of CAs. This creates problems of their own, e.g. modern web browsers have thousands of public key certificates belonging to different CAs preinstalled to be trusted. If only one of them is compromised or one CA does not do its job properly, this opens the window for attacks (Ferguson 2010, 284).²¹

A PKI is one of the best tools to solve the problem of key distribution in cryptography and it seems to be particularly suitable for our case, since we already have a centralized server, that could easily be used for key distribution. Nevertheless, we will have to take a close look at the details, since in our threat model we decided not to put (too much) trust in the server.

3.5.3 X.509 certificates

Certificates, of course, have been standardized and the most common standard are X.509 certificates (IETF 2008). An X.509 certificate consists at least of the following parts:

- Version Number
- Serial Number
- Signature Algorithm ID
- Issuer Name
- Validity period
- Subject name

²¹For anyone interested in this problem and a proposal of a possible solution we recommend the paper by Kim et al. (2013).

- Subject Public Key Info
- Public Key Algorithm
- Subject Public Key
- Certificate Signature Algorithm
- Certificate Signature

Depending on the version of the certificate there may be additional entries. The entries in the certificate make sure the subjects name (which can be any kind of identifier, e.g. an email address or a company name) is tied to the public key, while also making some indications on how the authentication works (Ertel 2012, 130f).

3.5.4 The Socialist Millionaires' Protocol (SMP)

The Socialist Millionaires' Protocol (SMP) is a zero-knowledge-protocol (see 3.4.1). In Off-the-Record Messaging (OTR) it is used for key verification in a synchronous communication session (OTR Development Team 2017). The premise is, that both parties share a secret. This secret can be relatively weak like a dictionary word, because the protocol takes place in real-time, so that an attacker has almost no time to perform a brute force attack on the secret.

In OTR once Alice and Bob have exchanged their public keys in an Authenticated Key Exchange (AKE)²² either party can ask for key verification at any time. The fundamental idea is, that Alice and Bob share a secret p . This secret can be something they have agreed on, in a (previous) out-of-band conversation or any knowledge only the two of them share. Note that it is not unlikely that such a shared secret exists in the context of instant messaging. Alice uses p to compute a secret value $x = h(P_A, P'_B, p)$ and Bob does the same computing $y = h(P'_A, P_B, p)$, where P_A P_B are Alice's and Bob's public keys and P'_A and P'_B are the public keys, they received from each other. This means $x = y$ is only true, if both $P_A = P'_A$ and $P_B = P'_B$ hold and they used the same secret p . The SMP allows Alice and Bob to check whether $x = y$ without leaking any other information to either party. If the statement is true, both know, that they share the same secret and use the same keys (Hallberg 2008).

Similar to Diffie-Hellman Key Exchange Alice and Bob agree on a 1536-bit prime number n and a number g , that is (at best) a primitive root modulo n . These values can be publicly known. All exponentiations in the following are computed modulo n :

1. Alice picks three random numbers $a, a', r \neq 0$ and computes g^a and $g^{a'}$.
2. Alice sends Bob g^a and $g^{a'}$.
3. Bob verifies $g^a, g^{a'} \neq 1$.

²²The AKE is based on the Diffie-Hellman-Key Exchange. For details see OTR Development Team 2017.

4. Bob picks three random numbers $b, b', s \neq 0$ and computes g^b and $g^{b'}$.
5. Bob computes $g' = (g^a)^b$, $g'' = (g^{a'})^{b'}$, $P_b = g''^s$ and $Q_b = g^s g'^y$.
6. Bob sends Alice g^b , $g^{b'}$, P_b and Q_b .
7. Alice verifies g^b , $g^{b'} \neq 1$.
8. Alice computes $g' = (g^b)^a$, $g'' = (g^{b'})^{a'}$, $P_a = g''^r$, $Q_a = g^r g'^x$ and $R_a = (Q_a Q_b^{-1})^{a'}$.
9. Alice sends Bob P_a , Q_a and R_a .
10. Bob computes $R_b = (Q_a Q_b^{-1})^{b'}$ and $R_{ab} = R_a^{b'}$.
11. Bob verifies that $R_{ab} = P_a P_b^{-1}$.
12. Bob sends Alice R_b .
13. Alice computes $R_{ab} = R_b^{a'}$.
14. Alice verifies $R_{ab} = P_a P_b^{-1}$.

Note that g' , g'' and R_{ab} are the keys negotiated in three Diffie-Hellman Key Exchanges. Therefore the following holds:

$$\begin{aligned}
R_{ab} &= (Q_a Q_b^{-1})^{a'b'} \\
&= (g^r g'^x (g^s g'^y)^{-1})^{a'b'} \\
&= (g^{r-s} g'^{x-y})^{a'b'} \\
&= (g^{r-s} g^{ab(x-y)})^{a'b'} \\
&= g^{a'b'(r-s)} g^{a'b'ab(x-y)} \\
&= g''^{r-s} g^{a'b'ab(x-y)} \\
&= g''^r (g''^s)^{-1} g^{a'b'ab(x-y)} \\
&= P_a P_b^{-1} g^{a'b'ab(x-y)}
\end{aligned}$$

$R_{ab} = P_a P_b^{-1}$ implies $g^{a'b'ab(x-y)} = 1$ and because $a, a', b, b' \neq 0$, $x - y = 0 \Leftrightarrow x = y$ holds.

Therefore the final verifications convince both parties, that the other party has knowledge of the same secret value, derived from both public keys and the shared secret, and the keys are successfully verified (OTR Development Team 2017).

Concluding the section on cryptography, we state that cryptography is one the most powerful technical tools to make a system secure and keep users' data private. Anyway, it is important to keep in mind that no cryptographic system can provide absolute security. A brute force attack always has a small chance to succeed and further pitfalls lie in implementation errors that lead to vulnerabilities of the whole system. Most importantly it is important that cryptography by itself does not render any system secure, but rather the security of a system depends on the correct

implementation of cryptographic methods. This includes the notion, that one of weakest links in any system is always the user. Therefore the system should be designed such that the user cannot unwillingly compromise the security of the system, while at the same time solving all the hard problems technically and making the system as intuitively usable as possible.

4 Related Work and similar projects

In a first section of this section we will discuss the research that has been done in the field of cryptographic cloud storages. In a second section existing software will be analyzed. The key question in both sections is how users can on the one hand securely sync and share data in cloud storage and how this security will not lead to sacrifice usability.

4.1 Research on cryptographic cloud storages

In recent years a lot of research has been done on cloud storages and how to enforce their security and privacy properties. In this section we will give a brief overview over the most relevant papers for the subject of this thesis. Some of them will be discussed in detail in the corresponding sections.

Grolimund et al. (2006) propose a solution for efficient and secure key management in Sync & Share applications such as cloud storages. Their cryptographic tree structure, called “Cryptree” is the structure implemented by most of today’s end-to-end encrypted cloud storages. Like us they introduce lazy re-encryption and separated read and write tree structures. However, their Cryptrees also contain so called “backlink keys” to access the metadata of parent folders. From our perspective this makes certain operations unnecessarily complicated and also represents a threat to user privacy, because a user might get information on data in parent folders that has not actually been shared with her. This is why we present a solution that works without backlink keys. We trust Grolimund et al. in their evaluation and believe that our own system will at least achieve the same performance and better in some cases due to the removed backlink keys.

Lám, Szebeni, and Buttyán (2012b) present Tresorium, which is the basis for the CSP Tresorit. They base their approach on Cryptree. In contrast to our design Tresorium does not allow to distinguish between read and write access when sharing data. Their approach for key management and distribution scales better when sharing data with groups, because their focus is on a mechanism to negotiate and distribute group keys. This mechanism is further elaborated in Lám, Szebeni, and Buttyán 2012a.

Zwattendorfer et al. (2013) resolve the problem of key distribution by introducing a PKI, based on the Austrian ID card. Unfortunately this requires both access to a Smart Card and trust in the CA of that PKI, the Austrian state, which is why we discard their solution for our application scenario. We analyze their approach in detail in section 7.3.

Körner and Walter (2014) also try to resolve the problem of public key distribution and verification in end-to-end encrypted cloud storages by using QR-Codes as a carrier medium. They encode the complete public key in a QR-Code. For us this solution is not considered usable enough, because it impedes sharing, when users have not physically met before, which is why we make leverage of the server for key distribution. Yet, inspired by their paper, we offer a QR-Code mechanism to verify other user’s certificates (see section 7.4).

Zhao et al. (2010) introduce a progressive encryption scheme which allows data to be encrypted multiple times by clients and be decrypted in a single operation. Unlike our design such a scheme allows multiple users to work on documents simultaneously. However, there are some limitations, the authors themselves mention. Most notably it creates a large computation overhead and it requires cooperation from the cloud provider that needs to share its own private key with all users a file is shared with.²³ Due to these limitations we chose not to implement their mechanism, although it provides interesting properties for collaborative and simultaneous work on documents.

Kamara and Lauter (2010) present a paper on cryptographic cloud storage, that uses modern functional cryptography to enforce access policies on the data stored and shared in the cloud. While these policies mostly aim at a corporate environment, the functional cryptography used by the authors also provides searchable encryption. However, the authors do not implement their proposed system, because certain obstacles still have to be overcome. While some of the properties presented are very interesting, more research has to be done on the functional encryption schemes before the solution becomes practically and efficiently feasible.

Wilson and Ateniese (2014) evaluate existing CSPs that claim to provide secure end-to-end encrypted cloud storage. Similar to our evaluation (section 4.2) they find that all existing CSPs lack certain security properties. In particular they point out that the CSP should not be trusted as the CA, if a PKI is used (cf. section 7.2). However, they only analyze a limited number of systems excluding for example Nextcloud. Since their paper is from 2014 they also cannot consider all recent changes in the evaluated systems, which is why we decided to do our own evaluation in the following section.

In recent years there have been different approaches to improve on the shortcomings of PKIs and to resolve the problem of end user public key verification in particular. One such approach is CONIKS (Melara et al. 2015), which allows users to monitor their key bindings and thus make sure no false public keys can be published. This removes the need for a trusted CA. However, a traditional PKI, in which the server acts as a CA and key distributor is better suited for the architecture of our system. Therefore, we do not implement CONIKS, which has not been widely adopted so far, but provide custom key verification by QR-Code scanning and the Socialist Millionaires’ Protocol.

In our literature research we came across many more papers that try to resolve the problem of end-to-end encrypted cloud storage (cf. Gadhavi et al. (2016), Rewagad and Pawar (2013), Kaur and Singh (2013) and Kumar et al. (2012)). These cannot all be presented in detail here, because they either aim at different security goals

²³Still, the cloud provider does not get access to the plaintext of the shared data.

than we do or do not provide any contributions, that we have not talked about in the before mentioned papers.

4.2 Comparison of existing client-side encrypted cloud storage systems

In this section we will compare existing cloud storage systems that implement client-side encryption. Most of these systems are not open source, which means we have to base our analysis on the descriptions they provide in their white papers and other publications. Since these sources are also used for marketing and often lack important technical details, they should be carefully judged. However, the goal of our comparison is not so much to judge how secure a certain system is in practice or to make recommendations about which systems to use, but rather to get an idea, how the tasks described in the previous sections can be tackled and how different mechanisms work together in existing encrypted cloud storages. The objective here is to analyze, which consequences each mechanism has for the functionality, security and usability of the system as a whole.

In our comparison we will focus on cloud storage systems that have a focus on private users and small-scale businesses. This means we will not explain in detail the solutions offered for the deployment in large- and medium-scale businesses,²⁴ but focus on systems that are available to any user. It also implies that our focus is not on the integration with existing (business) applications or scalability, but mostly as mentioned before on security, functionality and usability.

In our analysis we will therefore describe the functionality provided by a system, describe the cryptographic methods involved with a particular focus on key management, and summarize the pros and cons of the solutions regarding security and usability.

4.2.1 SpiderOak

SpiderOak has become famous as a Dropbox alternative recommended by Edward Snowden in 2013 (Kiss 2014). It then claimed to be “Zero-Knowledge” and later changed the term to “No-Knowledge”, since the system does provide end-to-end encryption, but the login mechanism does not fulfill the requirements of a cryptographic zero-knowledge-proof (SpiderOak 2017c).

Functionality: SpiderOak is a backup software. Its main objective is to provide a client (for all common Desktop operating systems) that automates the process of frequently uploading local backups to the cloud. The user can decide whether to backup just one file, some folders or the complete local hard drive. Due to its focus on backups SpiderOak does not allow to share files with other users, but it is possible to make files and even directories publicly available via a web link. While these links still need a login password to access the files, these files are not stored end-to-end encrypted on the server, but only with server-side encryption (SpiderOak 2017b).

²⁴Anyone interested in such systems might take a look at Syncplicity (2015) and FTAPI (T. 2014).

Security: SpiderOak uses hybrid encryption. The data is encrypted on the local machine before being uploaded to the SpiderOak servers. Each file and folder has its own key. Files and file metadata are encrypted using AES-256. Integrity and authenticity are ensured using HMAC-SHA-256. The keys are stored on the SpiderOak servers encrypted with the user password using PBKDF2 with 16000 rounds of SHA-256 and a 32 byte salt. SpiderOak calls this privacy environment “No-Knowledge” and claims that the “server sees only a sequentially numbered series of data containers” (SpiderOak 2017a). The fact, that the keys are stored on the servers alongside the data and only encrypted with the user password makes key exchange between devices a trivial problem. The user simply has to login providing her username and password, thus synchronizing keys with the existing account. SpiderOak also uses data deduplication²⁵ across the data of a single customer not across customers. While it would be possible to do this locally and detect duplicate files by creating an encrypted database of hash values of the plaintext of every file, SpiderOak unfortunately does not provide information on how data deduplication is actually done (SpiderOak 2017b, 2017a).

Usability and Conclusion: End-to-end encryption in SpiderOak is completely transparent, since the user only has to provide her username and password to access her account. SpiderOak is very easy to use from all common platforms and also has a web panel. The web panel, however, is only accessible after accepting a disclaimer, that informs about how the No-Knowledge policy is violated by accessing one’s account from the web panel.²⁶ Generally SpiderOak protects the confidentiality and integrity of the reliably stored data. However the security largely depends on the user’s choice of password. The major disadvantage SpiderOak suffers from, compared to other solutions, is that it is not designed to share the files stored in the application. Therefore its key management and key exchange mechanism do not provide much insight on how to an end-to-end encrypted sharing platform. As a backup storage SpiderOak provides a secure and easy-to-use solution.

4.2.2 Cryptomator

Cryptomator is not a cloud storage system but an application that is run locally and enables users to store their data protected by end-to-end encryption on a regular CSP’s servers (e.g. Dropbox, Google Drive, Nextcloud).²⁷

Functionality: While Cryptomator on its website claims to implement fully “transparent encryption”, it actually breaks the sharing function of cloud storage platforms it is based on. It only allows users to upload their encrypted files and access them, but it lacks a mechanism to exchange keys between users and thus makes it impossible to share encrypted files with other users. Although providing more or less

²⁵Data deduplication means that the exact same file is never uploaded and stored more than once in the cloud. A mechanism such as a comparison of hash values of files is used to detect duplicate files. Instead of uploading the duplicate, only a link is created to the file that is already stored on the servers. This allows to save bandwidth and storage.

²⁶The decryption in the web panel takes place in a JavaScript application, that is delivered by the server and should not be trusted.

²⁷There exist similar applications, like Boxcryptor, but since only Cryptomator is open-source we will only analyze Cryptomator in this paper.

the same functionality, the main difference to SpiderOak is, that Cryptomator creates a vault, that can be accessed just like an external drive such as a USB-stick (Cryptomator 2017a).

Security: Cryptomator realizes encryption similar to SpiderOak. Since it is only run locally, every encryption operation is performed locally before the data is stored in the cloud. Upon the creation of a vault, two 256 bit master keys are created for encryption/decryption (AES) and for authentication (HMAC). Upon encryption of a file or folder a unique AES-256 key is created and used for the encryption. This key in return is encrypted with the master key and stored alongside the encrypted file. The master keys are also encrypted with the user password using scrypt and stored in the root directory of the vault. To unlock the vault a user has to provide her password, which decrypts the two master keys and allows her to encrypt and decrypt files using the unique file encryption keys, that are stored in the vault. Cryptomator encrypts every file and its metadata that is added to the vault automatically before it is uploaded to the cloud, which is realized by the actual cloud storage application. It also decrypts every file automatically when accessed (Cryptomator 2017b).

Usability and Conclusion: Cryptomator provides good usability. Since everyone knows how to access a USB-drive, users will also be able to use Cryptomator. Data encrypted with Cryptomator is not accessible via a web interface. Due to the use of scrypt and two different keys for HMAC and AES-256, the integrity and confidentiality of the data is protected slightly stronger than with SpiderOak. However, a compromised password will always lead to a full compromise of the user's vault, for anyone with access to the user's data in the cloud storage. The lack of a mechanism to share files, will detain many potential users from switching to its services.

4.2.3 SafeMonk and Sookasa

SafeMonk was the first Dropbox Add-on that provided end-to-end encryption. In 2015 the developer, SafeNet, decided not to continue with the project and SafeMonk was merged into its former competitor Sookasa (Scherschel 2015). Since there is not much difference between both services, we will only describe SafeMonk.²⁸

Functionality: SafeMonk works on top of Dropbox, somewhat similar to the way Cryptomator does. However, there are certain differences in the key management, that allow for a better functionality. Like Cryptomator SafeMonk creates a vault, that is a special encrypted folder inside the Dropbox folder. The content of this folder will always be encrypted. SafeMonk does not only provide secure storage but also allows users to share data with other users, which means it basically provides the same functionality as Dropbox (Raymond 2015).

Security: In Safemonk all encryption, decryption and key generation operations are performed by the client. Unlike the solutions presented so far in SafeMonk the structure of the encrypted directories is reflected in the key structure. Every file has a unique encryption key (AES-256). This key is encrypted with the encryption key of the parent directory, which in return is encrypted with the encryption key from its own parent directory. This goes on to the root directory, the SafeMonk folder, which also possesses an encryption key, but this key is only encrypted with

²⁸There is a white paper for Sookasa (2017) available.

the user’s public RSA key. Organizing encryption in this manner creates a tree structure, where each node contains a key that decrypts its children’s keys. So every key can be restored to plaintext, if and only if the user owns the correct private RSA key to decrypt the root directory encryption key. The user’s private RSA key is stored on the server encrypted with the user’s password.²⁹ The tree structure has the purpose to allow users to share files and folders completely to their liking with no need for re-encryption, because a simple share of the folder key will allow the user to access all its content. This mechanism also makes sure other users only gain access to the files and folders they are supposed to access, while the rest of the data remains protected. Safemonk posts the user’s public key to the SafeMonk server. SafeMonk on the one hand claims that keys are stored in the client only (apart from the public keys) (SafeMonk 2013). However, it is possible to exchange keys between devices, by simply logging in to the Dropbox and SafeMonk accounts with the user’s passwords. It is also possible to access the encrypted data from the SafeMonk web interface with mere knowledge of the password. So, what they really do, is storing all user keys on their servers encrypted with the private key, which is also stored on the servers protected by the user’s password (Raymond 2015). The security of this mechanism is not completely trustworthy, because of this unclear information policy and, they do not state which password-based key derivation function is used to protect the private key.

In case Alice wants to share a file with Bob, she shares the file as usual using Dropbox. Upon receiving the share Bob, who apart from Dropbox also runs the SafeMonk client, requests the decryption key for that folder from the SafeMonk server. The Safemonk servers ask Alice to confirm the share with Bob. If Alice confirms it, her client will retrieve Bob’s public key from the SafeMonk server and use it to encrypt the folder key. This encrypted folder key is sent back to the server and from there to Bob, who now has access to the share (SafeMonk 2013).

Safemonk does not implement further measures to ensure the integrity and authenticity of the files stored in the cloud. Its main focus is on confidentiality only.

Usability and Conclusion: While the tree structure that allows very flexible and transparent sharing, is a great plus to the functionality of SafeMonk, it has some issues, that could be resolved better. First Safemonk uses its own servers for key exchange, thus creating a second infrastructure next to Dropbox, that could have been avoided with better integration. This infrastructure does not provide any additional security. In the described key exchange between Alice and Bob it is notable that the SafeMonk server could always act as a man-in-the-middle providing fake keys. This puts a lot of trust in the SafeMonk servers, since there is no mechanism to verify Bob’s public key. In our threat model this clearly is a break of data confidentiality. Although the key exchange mechanism is not secure, it still requires some simultaneous interaction from the users without a gain in security.

In general it can be stated, that the tree structure used for key management is very useful, while the mechanisms for key exchange are neither completely secure nor provide good usability. Another problem is that SafeMonk does not always explain how things are done, e.g. where the keys are really stored. Therefore SafeMonk should not be trusted entirely. Especially, it is dubious whether under

²⁹Unfortunately we did not find out, which algorithm is used to encrypt the private key.

these circumstances the code delivered by the web interface can be trusted not to compromise the user's password and secret keys.

4.2.4 Sync

Sync is a commercial CSP that focuses on privacy (at least in their marketing): “‘Zero-knowledge’ means we absolutely cannot access the encrypted data stored on our servers. Your data is completely safeguarded from unauthorized access, which is the only way you can completely trust the cloud. Protecting your right to privacy is our passion” (Sync 2017a).

Functionality: Sync provides functionality similar to that of Dropbox. It is possible to use Sync to backup files, synchronize files between devices or share them with others. It supports all common operating systems and also is accessible via a web panel. Sync also allows users to share data via web links.

Security: Sync's security is based on an RSA-2048 key pair, that is used to decrypt all other keys. The private RSA key is stored on the Sync server encrypted with the user password using PBKDF2. The initial generation of the key pair can be done by the client or inside the web panel. Although Sync claims to be “Zero-Knowledge” it sends a BCrypt hash value of the user password and a salt (fetched from the server) to the server in order to authenticate the user and be able to access the secret keys. BCrypt is a one-way cryptographic hash function and therefore does not fulfill the criteria of a zero-knowledge-proof.³⁰ In Sync each file is encrypted with a unique AES-256 key that is stored on the server encrypted with the user's private RSA-key. The metadata is encrypted with AES-256 using a key that is derived from the user's password with PBKDF2. When a user creates a share, a unique share key is generated and encrypted with public RSA keys of all users the data is shared with. This share key is used to encrypt all file keys (Key Wrapping). If a user creates a secure link to share files, the shared files are encrypted with the share password using PBKDF2 and AES-256 and the password is added to the URL after a hash tag (fragment identifier), so that it is never transmitted to the server, but can be used locally by the JavaScript code to derive the decryption key to unlock the share.

Sync can be accessed via a web panel. In this case JavaScript Code is delivered by the server and used for key generation and all other cryptographic operations. Sync is aware that this means putting a certain level of trust in the server and has thus made the JavaScript code for the web panel publicly available, so everyone can verify that no secrets (password or keys) are leaked to the server (Sync 2017b). However, when the actual code is delivered, no integrity check is performed in the browser, so that the publication mostly must be seen as a demonstration of good will by Sync, since it provides no actual protection.

Usability and Conclusion: Sync provides good usability, since the encryption is completely transparent. It also uses modern cryptography and can therefore in general be considered secure. Sync has the same problem like all services that rely on the user's choice of a good password. Because a weak password will make it

³⁰The server could easily confirm an assumption on what the user's password is by calculating the cryptographic hash value for that password.

relatively easy for the attacker to gain access to the whole account. Although the key management used in Sync allows users to share data flexibly, a lot of keys need to be wrapped each time a folder is shared. This task can be sped up by a cryptographic tree structure, that would also save bandwidth. One of the biggest flaws Sync has security-wise is the lack of a mechanism to authenticate the users, that the data is shared with. There is no way to verify other users' public keys and one has to completely trust the server in providing the correct keys.

4.2.5 Tresorit

Tresorit is similar to Sync, but takes some extra measures to ensure security. It not only protects the confidentiality of the data users upload to the cloud, but also focuses on protecting the integrity.

Functionality: Like Sync Tresorit's main services are file synchronization across devices and sharing data with other users. This includes data backups. Tresorit also allows users to access all their data from a web panel and share files by link.

Security: In Tresorit all cryptographic operations are performed on the client side. Each user has two pairs of private RSA-4096 keys, one for encryption (wrapping the symmetric AES-256 root key) and one for authentication and integrity checks (based on HMAC). These private keys are stored in the user's root directory on the server encrypted with the user's password using PBKDF2 with HMAC-SHA1. Before uploading each file is encrypted with a unique AES-256 key. To protect the integrity of the files Tresorit uses HMAC-SHA-512. Each upload and modification of a file is also authenticated with the user's digital signature using an RSA-4096 key and SHA-512 as a hash function. To store the keys Tresorit uses a tree structure similar to that of SafeMonk, i.e. each file key is encrypted with the key of the containing folder. Unlike Safemonk, Tresorit stores the keys alongside the files itself in the same folder and uses no extra infrastructure for key management. Using this key structure, sharing is realized by making the decryption key of the root directory of the share available to the users the directory is shared with. When sharing with a group of users there is one pre-master secret for each user granted access to the files. To unlock a share the user's private RSA key is used to decrypt this user's pre-master secret. The resulting key is used to calculate the AES-256 key of the shared directory (which unlocks the root key for that share) by computing the HMAC of the sharing users' public key certificates.³¹

The certificates used are issued by Tresorit itself and are X.509 certificates, that contain the user's email address as an identifier. This means Tresorit has its own PKI, in which it acts as the CA. However the authentication mechanism for another user's key is more sophisticated than in a standard PKI: When Alice wants to share a file with Bob, she invites him to the share by providing Bob's mail address. Tresorit then sends an email to Bob, that contains a 256 bit random secret. When Bob accepts the invitation in Tresorit, he has to provide this secret in a challenge-response protocol. Tresorit sends both, Alice and Bob, the certificate and the public

³¹Note, that there is an alternative version, which uses Diffie-Hellman certificates instead of pre-master secrets (Tresorit 2017b). The group key negotiation Tresorit used has been described in detail in Lám, Szebeni, and Buttyán 2012a.

key of their counterpart, which is supposed to guarantee that they really have the correct certificate for the user identified by the email address. However, if Alice is not certain, that this is secure enough, she can use an out-of-band channel to agree on a password with Bob and both will have to provide this password in order to verify the public key. Unfortunately Tresorit does not offer more details on how the secret and how the password are used to ensure the authenticity of the counterpart's public key. Cryptographic methods should use well proven algorithms to achieve security. Since Tresorit does not provide such a mechanism its key verification mechanism is not trustworthy. Tresorit allows users to revoke shares. In such case Tresorit deletes the user from the Access-Control-List (ACL) of this share and re-encrypts the root directory. This immediately impedes access to the root directory by the revoked user. In order to save computing power, the rest of the data is only re-encrypted, as it is changed (lazy re-encryption). This makes sure a revoked user will see no further changes made to a document she could formerly access. When a user logs in, a challenge-response protocol is applied. The server sends the user a challenge and the 160 bit salt. The user uses PBKDF2 to compute a key from the password and the salt. She calculates the HMAC of the challenge with that key. The result is the response which she sends back to server. The server has stored the key derived from the password and salt and also computes the HMAC of the challenge. If the result matches the response, the user is granted access to her account (Tresorit 2017b, 2017a).

Usability and Conclusion: Tresorit provides all basic cloud storage functionality like file synchronization, sharing and access via web panel and clients on all common platforms. It is the most secure of the presented solutions. This is due to its application of modern cryptography and its neat key management. It also has the most elaborated key exchange and is the only solution that puts a lot of effort in ensuring data integrity (using separate keys for encryption and integrity checks). However, its security is based on the user's choice of password, as the password unlocks everything, but users usually choose weak passwords. Also, when using Tresorit's web panel the user has to trust the server not to deliver malicious JavaScript code, that is used for the cryptographic operations. But since it is closed-source software users have to put the same trust in the client software. One major problem with Tresorit is, that it does not use a well-proven key verification mechanism, which means a Man in the Middle attack by the server (which acts as a CA) might always be possible (cf. 7.2).

4.2.6 Nextcloud

The ownCloud fork Nextcloud presented its design for end-to-end encryption and a proof-of-work Android client in September 2017. It is the first cloud software, that is open-source and therefore self-hostable, that provides end-to-end encryption.³²

Functionality: Nextcloud is a cloud storage software that can be self-hosted or deployed in a public cloud. It provides the basic sync and share functionality,

³²Seafile is also open-source software and generally performs encryption operations on the client, but for some use cases the password is sent to server and used to derive the key (Seafile 2017), which is why we do not consider it in this section.

i.e. users can upload their files, sync them between devices and share files with other users and user groups. It also offer a great number of additional applications such as calendars, password managers, audio/ video calls etc. (Nextcloud 2017b). Nextcloud after the integration of end-to-end encryption will only support end-to-end encryption for new empty directories. It means while encrypted and not encrypted folders may live in the same account, an existing unencrypted folder cannot be encrypted. End-to-end encrypted folders are not displayed and cannot be accessed by the web interface or older versions of the client that do not support end-to-end encryption.

Security: Nextcloud allows users to create encrypted folders. These folders can be shared with other users. For the encryption Nextcloud uses AES-128 with a unique key for each file. It encrypts all metadata (in a JSON object) and all file keys with an AES-128 share key, thus not giving away any information apart from the number of files contained in the folder. The AES encryption key is encrypted with the public keys of all users, that are supposed to get access. When setting up her account for end-to-end encryption the user generates an RSA key pair and retrieves a X.509 certificate from the server, that acts as a CA in the PKI used for key verification. The user's private key is encrypted with PBKDF2 with HMAC-SHA1 using 12 randomly chosen words from a 2048 words list as a password. When the user sets up the client on a new device, it downloads the public key certificate for her user ID and the corresponding encrypted private key. If they match, the user can decrypt the private key with the 12 word password and store it on the device. When a user wants to share a folder with other users, she first has to get their certificates from the server and make sure they are signed by the the CA, the server. When a user gets revoked from an existing share, lazy re-encryption is applied (Nextcloud 2017a).

Usability and Conclusion: The best thing about Nextcloud is that it is the first open-source cloud storage software that provides client-side encryption. However, there are some pitfalls. Regarding usability the most significant disadvantage is probably that Nextcloud is inflexible with sharing. Users need to create special folders for encryption and cannot just encrypt and share existing folders. Also it is impossible to re-share a subfolder of an existing encrypted share. Nextcloud could have benefited a lot from a cryptographic tree structure. However it provides a secure and user friendly way for cross-device key exchange. Momentarily key verification is based completely on the trust in the CA and users are not allowed to change their private keys, which means that there is no way to recover from a compromised private key. Generally the end-to-end encryption in Nextcloud is a good approach, that is suitable for many use cases, but it lacks flexibility and proper key verification and key revocation mechanisms. Client-side encryption is not planned to be supported by the web panel due to the security risks this represents.

Summarizing our analysis of today's client-side encrypted cloud storage solutions it is worth pointing out that there is a number of systems available that uses secure cryptographic algorithms such as AES and RSA/ Elliptic Curves. These systems combine symmetric and asymmetric cryptography to achieve secure and efficient systems. However, most of these systems suffer from severe limitations, e.g. Spi-

deroak, Cryptomator do not allow users to share their data easily. Systems such as Sync, Nextcloud and Tresorit rely on a service provider managed PKI to provide the sharing functionality. Yet, these systems suffer from the lack of a secure certificate verification mechanism, that will make man in the middle attacks impossible. This is the case for Tresorit, which also suffers from weak protection of user’s private key. Nextcloud, on the other hand resolves this problem by forcing a 12 word random generated passphrase to protect the user’s private key. Nextcloud’s PKI, however, does not provide essential mechanisms such as certificate revocation in case of a private key compromise. In the following sections we will present our own systems that provides new mechanisms and combines the so far presented mechanism in a new way in order to provide a secure system, even vulnerable users can have full trust in when sharing their most confidential data.

5 Key Management

In this section we will propose a cloud storage architecture that allows to protect all the major security objectives for our threat model: confidentiality, integrity and authenticity of the stored data. At the same time it has to be flexible and must not break the core share and sync functionality of cloud storages, since otherwise users will simply stay with the insecure existing cloud storages. we will, however, not describe any details of basic functionality of cloud storages, but only focus on the security relevant aspects in the client, that is mainly the key management. It is obvious that we can apply the best cryptographic methods and not design a secure regarding our threat model and usable system at all.

In server-side encryption the cloud storage provider always has access to the keys. So the crucial question of key management is who has access to the keys. This thesis proposes a design that allows only the user, i.e. the owner of the data and other users that have explicitly been granted access, to access her encryption keys. This is what we call client-side encryption. Apart from the question, who can access the encryption keys, it is important to organize encryption in a way, that allows the system to be effective and does not sacrifice performance. Also, it should conserve the original functionality and usability as best as possible.

5.1 Architecture

In this section we describe the key management, that is how encryption is organized, how keys depend on one another, where they are stored and how operations are performed.

We propose a cryptographic tree structure, that has been used in practice in systems like SafeMonk and Tresorit. The design was first published by Grolimund et al. (2006) under the name “Cryptree” and our design is mainly an adaptation of Cryptree. Its main benefits are simplicity, efficiency and intuitive semantics. Note that there exist different variants of Cryptree, e.g. the inventors of Tresorit describe an early version of how they implemented it in their system “Tresorium” (Lám, Szébeni, and Buttyán 2012b) and Mathieu (Bourrier 2016) uses it in his draft of client-side encryption for ownCloud under the working title “ownCrypt”. Our own architecture

is greatly inspired by those approaches. However there are some disadvantages to both Cryptree and Tresorium. Cryptree overcomplicates key management by introducing a back link key structure for “Upward Inheritance of Access Rights”. This allows users to access the metadata of ancestor folders and turns the cleaning operation and the revocation operations into rather complicated procedures (see 5.2). Tresorium on the other hand does not allow to differentiate between read and write access and provides only one “tresor”, a shared folder, per group. The ownCrypt draft does not foresee to use lazy re-encryption and is therefore not optimal for frequent changes of the users granted access. The architecture proposed here is designed to provide maximum flexibility and performance, while completely conserving the confidentiality, integrity and authenticity of the data.

Since we want users to interact with the system as intuitively as possible leveraging the folder structure is the obvious choice. With the proposed design, we will get the following characteristics:

1. Sharing a folder also grants access to all subfolders and files contained in it.
2. Sharing a folder does not grant access to any ancestor folders or sibling folders (that are contained in the same parent folder). This includes metadata.³³
3. It is possible to differentiate between read and write access, where write access always includes read access but not vice versa.
4. File metadata is treated as confidential. It is considered acceptable to give away information on the size and number of files, but not other metadata (names, time of modification etc.) and the content.

To achieve this we will use key wrapping or as Grolimund et al. (2006) put it: “cryptographic links”. This includes both symmetric and asymmetric links. To create a cryptographic link from $K_1 \rightarrow K_2$, K_2 is encrypted with K_1 , thus anyone who has access to K_1 and the link, also has access to K_2 . Cryptographic links can be interpreted as the edges of a graph, where the keys themselves are the nodes. For every folder there is one read item and one write item. For every file there is one read item, one write item and the file ciphertext which is named with a random identifier. In these items the keys and the data needed to organize encryption are stored.

In order to differentiate between read and write access, two different cryptrees are needed. First we describe the read access cryptree as shown in Figure 1. Every folder has a unique read key and a read item, in which the following information is stored:

- the metadata (name, creation, modification time, size etc.) of the folder (encrypted with the folder’s read key)

³³This is different from what Grolimund et al. (2006) proposed, since we do not want users to have upward inheritance of access rights which would allow them to get information on the ancestor folders of a folder or file that has been shared with them.

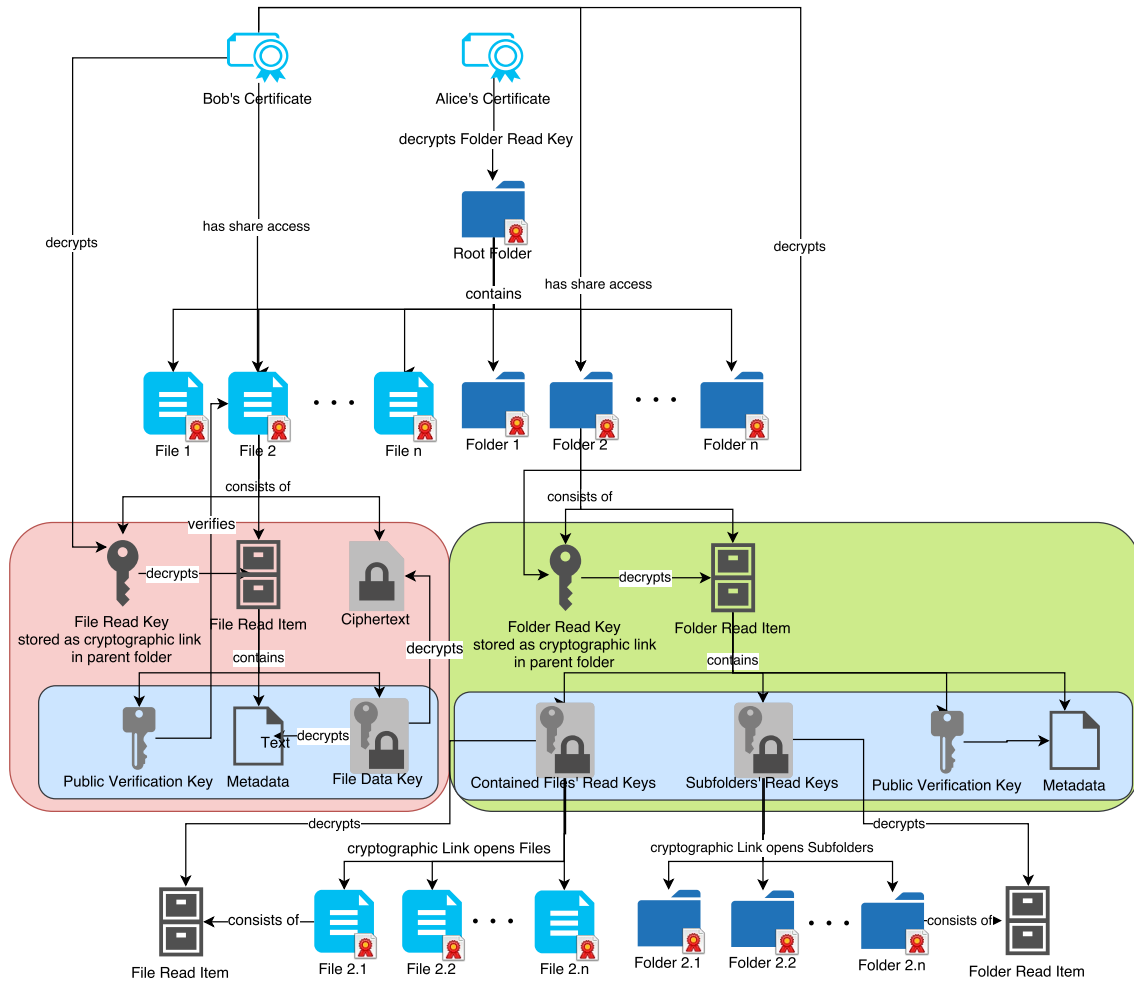


Figure 1: An example of a Read Crypttree

Red box: File 2. Green box: Folder 2. Blue boxes: the respective file/folder items.

Note the asymmetric cryptographic links from Bob's certificate to the File/Folder Read Keys are actually stored in the corresponding File/Folder Read Item (but of course not encrypted with the File/Folder Read Key).

- a cryptographic link with a unique subfolder read key for each subfolder that is used to encrypt the subfolder's metadata and create further cryptographic links
- a cryptographic link with a unique read key for each file that is used to encrypt the file data key
- a (public) verification key to verify the signature of all operations and the authenticity and integrity of this verification key itself
- an asymmetric cryptographic link for each user with read access to the folder from her public key to the folder key (only needed for the root folder and shared folders, not for the subfolders, that are implicitly shared)

For each file there is the file's ciphertext a read key and a read item, in which the following information is stored:

- the file data key used to encrypt the file as stored on the server (encrypted with the file's read key)
- the metadata (name, dirty flag, modification time etc.) of the file (encrypted with the file's data key)
- a public verification key to verify the signature of all operations and the authenticity and integrity of this verification key itself
- an asymmetric cryptographic link for each user with read access to the file from her public key to the file's read key (only needed, if the file is explicitly shared with that user)

For write access there exists a separate tree structure that is similar (see Figure 2). Public key cryptography is used to check the authenticity and integrity of write operations. The idea is that every change of the contents or metadata of files and folders has to be signed and only those with write access have the correct signature keys, but everyone with read access to a folder can check whether the signature is valid, because the verification keys are stored in read folder and file items. If not, the file is rolled back to the last valid version.

For each folder there is a write key and a write item, that contains the following information:

- a private signature key for all write operations (corresponding to the public verification key in the folder's read item) encrypted with the folder's write key
- a symmetric cryptographic link for each subfolder from the write key to the subfolder's write key
- a symmetric cryptographic link for each file from the write key to the file's write key
- an asymmetric cryptographic link for each user with write access to the folder from her public key to the folder's write key (only needed for the root folder and shared folders, not for the subfolders, that are implicitly shared)

For each file there is a write key and a write item that contains the following information:

- a private signature key for all write operations (corresponding to the public verification key in the file's read item) encrypted with the file's write key
- an asymmetric cryptographic link for each user with write access to the file from her public key to the file's write key (only needed, if the file is explicitly shared with that user)

This structure allows granting read access independently from write access. However since a lot of signatures have to be checked, handling write access separately from read access can slow down the system significantly for certain operations (e.g. the revocation of write access requires a lot of large public/private key pairs to be re-generated and a lot of write items have to be re-encrypted because lazy revocation

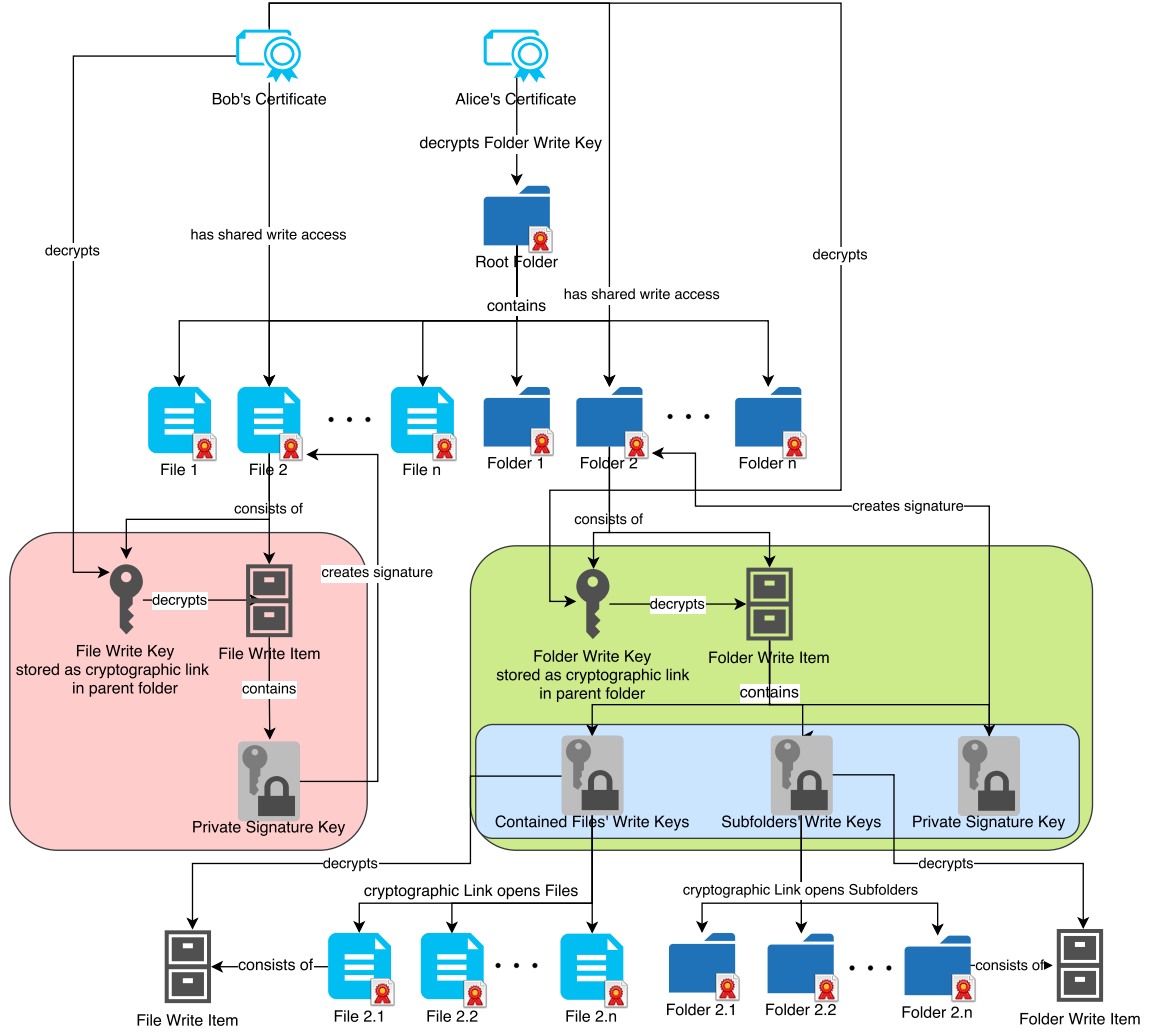


Figure 2: An example of a Write Crypttree

Red box: File 2. Green box: Folder 2. Blue box: the respective folder item.

Note the asymmetric cryptographic links from Bob's certificate to the File/Folder Read Keys are actually stored in the corresponding File/Folder Read Item (but of course not encrypted with the File/Folder Read Key).

is not applicable here. So one should consider carefully whether the separation of read and write access is really needed or whether everyone with read access should automatically have write access which would cause a boost to performance of the system, since the write item tree could be left out completely by adding the signature keys to the read items.

To achieve an acceptable performance our design has to avoid asymmetric cryptography, where possible. This is realized by giving each user only one private/public key pair, that serves as her identification and grants access to the keys used for data encryption. Because we use a tree structure that only consists of symmetric keys, the amount of public key cryptography needed is marginal. However, some additional public-key cryptography is needed to make sure the data can be checked for integrity and authenticity. Public-key cryptography in this scenario cannot be avoided since we do not only want users with write access to be able to perform the

checks, but all users with read access.

Apart from avoiding public-key cryptography it is particularly important to avoid frequent re-encryption and changes of the encryption keys. The tree structure allows us to re-encrypt only those files that were actually changed reducing the use of bandwidth to a minimum because we do not have to re-encrypt a complete folder and upload it again when a single file in it changes.

The proposed design guarantees the confidentiality of the user's data when she revokes access for another user, i.e. this user loses access. Considering very dynamic user groups, i.e. with frequent changes of the group members, it is not feasible to substitute all the involved keys and re-encrypt all data, as this might cause an enormous use of bandwidth (and entropy and computing operations of the client). The application of lazy re-encryption resolves this issue. Lazy re-encryption means, that when a user is revoked access she immediately loses access to the encryption keys because the cryptographic link using her public key is immediately removed from the share, but if the user has access to the server and prior to the revocation made a local copy of the encryption keys, she can still access the files. This is considered acceptable, because anyone with a local copy of encryption keys will probably also have a copy of the actual data. So no additional information is leaked. However, lazy re-encryption ensures that a user who has been revoked access to the share cannot keep track of any further changes. Therefore once a shared files changes, a new encryption key is generated and used to re-encrypt the new version of the file. Lazy re-encryption in our design is realized by marking files as dirty, when someone is revoked access and cleaning them before changes are uploaded. Because the system also aims to support, that users without write access to parent folders are able to clean the files contained in a folder, folder keys have to be kept clean at all times and an additional data key is needed in the file item. File data keys are the only keys that can ever be marked dirty, because all other keys have to be renewed at once. So even with our realization of lazy re-encryption some folder and file keys have to be re-encrypted immediately, but the system avoids the re-encryption of files, that are usually much larger than 256 bit encryption keys. This saves a significant amount of bandwidth. Unfortunately lazy re-encryption cannot be applied to write access rights, because nobody could distinguish between legitimate and illegitimate signatures. Therefore verification and signature keys are renewed immediately at the revocation of write access. This makes write access revocation one of the most costly operations in the proposed design.

5.2 Operations

In this section we will only describe the operations concerning key management and cryptography but not the operations that need to be done in the underlying cloud storage infrastructure. We will not explain trivial operations, that just work straight forward or can easily be composed of other basic operations.³⁴ Some operations require certain access rights, this is always assumed to be the case. For example moving a folder requires write access to both the old and the new parent folder. In the proposed system granting and revoking access rights can be realized by anyone

³⁴E.g. a file/folder is replaced by removing the old file/folder and adding the new one.

with write access to the shared folder/file.³⁵

Every operation is performed in the client and handled as a transaction, i.e. it is either completely finished (commit) or not at all (rollback). The system does not support users to work on the files in parallel, because the server does not have any information on the files' content and can therefore not make use of versioning.³⁶

5.2.1 Adding a file

The client generates a new read key, write key, data key and a signature and verification key pair for the file. The file's read item is created containing the metadata, the data key and the verification key using the read key for encryption where it applies. The read item is named with a random identifier. The file is encrypted with the data key and the ciphertext (also named with the file identifier) and the read item are stored in the cloud. A cryptographic link to that new read key is stored in the parent folder's read item. The file's write item is created containing the signature key encrypted with the write key. The write item is stored in the cloud using the same identifier as the ciphertext and the read item. In the parent folder's write item a cryptographic link to the file's write key is created.

5.2.2 Removing a file

To remove a file its read and write items and ciphertext of the file are deleted. The cryptographic links are removed from the parent folder's read and write items.

5.2.3 Moving a file

Cryptographic links in the new parent folder's read and write items are created pointing to the read key and to the write key of the file respectively. The cryptographic links pointing to the read key and the write key are removed from the old parent folder's read and write items. The file is marked dirty.³⁷

5.2.4 Copying a file

First a physical copy of the file, its keys and read and write items is created in the old parent folder. Then this copy is moved to the new parent directory.

³⁵Note that this means that it is theoretically possible to deny the owner access to its files. However, apart from what is cryptographically possible there will be an ACL to manage the users privileges. So the owner might choose a configuration that restricts operations, that will render her unable to access her own data.

³⁶It might be feasible to perform versioning by clients with write access by merging the changes made in parallel versions and then uploading a document that includes all of them. Note that this would only be possible for certain file types.

³⁷This makes sure that everyone with access to that particular file keeps it, while those with access to its old parent directory lose access. Everyone with access to the new parent folder gains access.

5.2.5 Updating a file

The file's read and write items are cleaned (see below). To update a file the file data key is used to encrypt the new plaintext and the old ciphertext is replaced with it. The new metadata is encrypted with the file data key and the entry updated in the file's read item.

5.2.6 Adding a folder

The client generates a new read key, write key and a signature and verification key pair for the folder. The folder's read item is created containing the metadata and the verification key using the read key for encryption where it applies. The folder's read item is named with a random identifier and stored in the cloud. A cryptographic link to that new read key is stored in the parent folder's read item. The file's write item is created containing the signature key encrypted with the write key. The write item is stored in the cloud using the same identifier as the read item. In the parent folder's write item a cryptographic link to the file's write key is created. Subfolders and contained files are added recursively.

5.2.7 Removing a folder

All subfolders and files are removed recursively from the folder. The folder's read and write items are removed. The cryptographic links are removed from the parent folders read and write item.

5.2.8 Moving a folder

Cryptographic links in the new parent folder's read and write items are created pointing to the folder's read and the write key respectively. The cryptographic links pointing to the folder's read key and the write key are removed from the old parent folder's read and write items. The folder is marked dirty.³⁸

5.2.9 Updating a folder

To update a folder (e.g. renaming it) the new metadata is encrypted with the folder's read key and the metadata entry in the read item is replaced with it.³⁹

5.2.10 Granting read access (Sharing)

To grant read access to a folder/file to another user a cryptographic link from the user's public key to the read key is added to the read item of the folder/file.

³⁸This makes sure that everyone with access to that particular folder keeps it, while those with access to its old parent directory lose access. Everyone with access to the new parent folder gains access.

³⁹Note that folders actually only consist of metadata. So there is no need to apply actions recursively.

5.2.11 Granting write access (Sharing)

To grant write access to a folder/file or file to another user a cryptographic link from the user's public key to the write key is added to the write item of the folder/file. The user is granted read access to the folder/file.

5.2.12 Revoking write access

When a user is revoked write access to a folder/file, the cryptographic link from her public key to the write key is removed from the write item. The write item of the folder/file is cleaned.

5.2.13 Revoking read access

The user's cryptographic link is removed from both the folder/file read and the write item and the folder/file is marked as dirty.⁴⁰

5.2.14 Mark a file as dirty

The file's write item is cleaned. The dirty flag is set in the metadata in the read item. A new read key is generated and the cryptographic link in the parent folder's read item is updated to point to that read key. The new read key is used to re-encrypt the data key in the read item of the file. The asymmetric cryptographic links in the read item are updated to point to the new read key.

5.2.15 Mark a folder as dirty

The folder's write item is cleaned. A new folder key is generated and the cryptographic link in the parent folder's read item is updated with that folder key. The new folder key is used to update the cryptographic links to the subfolders and contained files in the read item of the folder. The asymmetric cryptographic links in the read item are updated to point to the new folder key.⁴¹

5.2.16 Clean a file read item

If the file is not marked dirty, nothing needs to be done. Otherwise a new file data key is generated and used to re-encrypt the file and its metadata in the file read item. The file data key is updated in the read item using the read key to re-encrypt it. The dirty flag is removed from the file's read item.

5.2.17 Clean a folder read item

Folders are always kept clean (see 5.2.15) and the contained files can be cleaned individually when needed.

⁴⁰Revoking read access includes revoking write access.

⁴¹All keys are immediately updated, because otherwise updating a dirty folder would require access to the parent folder in order to update the cryptographic links in it. This would make it impossible to clean folders for users that do not have full access. For this reason we also use separate read keys and data keys for files.

5.2.18 Clean file write item

A new signature/verification key pair and write key are generated. The file's write item is updated storing the new signature key encrypted with the new write key and updating all cryptographic links for each user using the new write key. The cryptographic link to the write key for the file is updated in the parent folder's write item. The verification key is updated in the file's read item.

5.2.19 Clean folder write item

A new signature/verification key pair and write key are generated. The folder's write item is updated storing the new signature key encrypted with the new write key and updating all cryptographic links for each user using the new write key. The cryptographic link to the write key for the folder is updated in the parent folder. The verification key is updated in the folder's read item. All write items of subfolders and contained files are cleaned recursively.

In this section we described an architecture the most important file and folder operations for the key management in our system. Our design allows for flexible sharing and revocation. Contrary to existing the before mentioned existing solutions our design not only allows to distinguish between read and write access but it ensures cryptographically that users with share access to a folder cannot access any meta-data of the parent folders. Most of the file and folder operations are similar to Grolmund et al. (2006), which is why we relinquish to do an evaluation of the cryptographic tree structure and instead focus on our key exchange and key verification mechanism presented in the following two sections.

6 Cross-device Key Exchange

In this section we will design the mechanism to exchange private/public key pairs between a user's devices. We will mostly follow the protocol described in (Nextcloud 2017a). Prior to explaining any key exchange mechanisms, it is worth pointing out that in our design, there will only be one private/public key pair (with the corresponding X.509 certificate) for each user. The alternative to this design would have been to use a unique private key for each of the user's devices. This would have made the revocation of a device easier, because only one device would have been needed to be revoked. Also, it would have allowed users to define for every file and folder, which devices are granted access.⁴² However, usually users want to use the cloud storage to sync all their data between all their devices, so this feature might easily be turned into a major inconvenience causing a lot of micromanagement of devices. A user that wants to give another user access to a folder cannot be bothered to decide which of the other user's devices should be granted access. So, in order to make it easy enough for users to share data, we would need to introduce a master private/public key, that is used for all shares with its owner. But this brings up yet another question: Which devices should have access to the private master key? As

⁴²Such a mechanism is planned for ownCloud (Bourrier 2016).

we can see, such a design would have led to a very complicated key management and sacrificed some of the intuitive usability of our current approach, because it is simply too complicated for the average user. The benefits (such as device revocation and access management per device) are relatively small and only apply to some special use cases. This is why we decided that each user only owns one private/public key pair.⁴³

In order to grant access to all of her devices the user only needs to transfer this key pair to these devices. To realize this task, we are going to use the server for key distribution, both public key certificates and encrypted private keys. More precisely we use a PKI, in which the server acts as a CA. This will allow us to distribute keys rather easily and also generate some basic trust, because it means in order to provide a fake public key or get access to a user's encrypted private key an attacker needs access to the server. Since we do not want to trust the server completely, further mechanisms for key verification are explained in section 7.

6.1 Adding the initial device

When the client first registers for encryption it generates an RSA-4096 key pair. The client requests a X.509 certificate from the CA, the server, using the current username as an identifier and the just generated public key. If there is no existing certificate for the clients username and the requesting client's username matches the provided username the server issues the certificate, stores it and sends it back to the client, i.e. the principle of Trust On First Use (TOFU) is applied. The client stores both the private key and the new certificate securely on the device. The X.509 certificate is publicly accessible on the server for any registered client.

In order to be able to recover from a lost device and to transfer the keys to another device, in a second step the private key is encrypted and stored on the server: the client randomly generates a twelve word passphrase from the BIP-0039 word list, which contains 2048 words. This means there are 2^{23} uniformly distributed passphrases (given that strong random numbers are available). The client derives a key from this passphrase using PBKDF2 with HMAC-SHA256 and encrypts the private key with it, generating a X.509 private key. This encrypted private key is stored on the server, accessible only for a client with the current username. The passphrase is displayed on the client and stored securely on the device.⁴⁴ It is never transmitted.

6.2 Adding further devices

When the client registers a new device, it checks whether there is an existing X.509 certificate for the username. If not, see 6.1. Else, the client retrieves the certificate from the server and verifies that the username in the certificate matches the one the client uses. Next the client downloads the corresponding X.509 private key from the server and asks the user to provide the passphrase to decrypt it. The client checks

⁴³Note that this key pair can get revoked and replaced by a new one.

⁴⁴How cryptographic keys and passphrases are stored on a device depends on the used platform. Usually the operating system offers a key chain for that purpose.

the decrypted private key really corresponds to the X.509 certificate it received. If so, it stores the private key, the passphrase and the X.509 certificate securely on the device.

Note this key exchange mechanism protects 1) confidentiality, because only a user with knowledge of the passphrase is able access the plaintext of the private key, 2) integrity, because the client checks whether the private key matches the X.509 certificate and 3) authenticity, because the received private key and corresponding X.509 certificate have been stored by someone with knowledge of the user passphrase, i.e. a client running on one of the user's devices. We chose to use the server to facilitate the cross-device key exchange, because this way the user does not have to transfer keys manually, which is neither secure nor very user-friendly. On the other hand most of the client-side encrypted cloud storages, we presented earlier, let the users choose their own passwords. Since users rarely choose strong unique passwords, we force them to use a strong randomly generated passphrase. Since the passphrase consists of English language words, it is relatively easy to type. The effort required from the user is acceptable, because the user has to provide the passphrase only once for each new device. This effort is repayed by the much better security it provides compared to letting users choose their own passwords as in other solutions such as Tresorit and Sync.

6.3 Revoking a X.509 certificate

Nextcloud does not allow users to change their X.509 certificates and private keys. While it may be true, that users can recover from a lost device, because they have left a note with the passphrase in a secure place or have more devices that still have access to the private key (and passphrase), Nextcloud completely overlooks the threat imposed by a compromised private key (or passphrase). Although the key is stored securely on the device and never transmitted in plaintext, a key compromise may still occur, e.g. if the user loses her device and only has a weak password in place to protect the device's key ring. In such case Nextcloud does not provide any mechanism to revoke access for a lost device and recover from the compromise, because a user can never change her private key. Nextcloud argues enabling users to renew their certificates and public keys, would open the window for attacks, because an attacker could impersonate the user and provide a new certificate to the server. This threat is similar to the man-in-the-middle attack discussed in section 7.2. In Nextcloud, to make such an attack less likely, a user will only once retrieve another user's X.509 certificate from the server and then store it locally. In all following shares it will use the same certificate. Because this only gives the server one opportunity to provide a fake certificate, this principle is called Trust On First Use (TOFU).

The ability to revoke a X.509 certificate is considered a fundamental security feature. Also, in our design users do not exclusively base trust in the certificates they retrieve from the server, but have additional mechanisms to verify each other's certificates. Therefore renewing a private/public key pair is acceptable.⁴⁵

⁴⁵Tresorit (2017b) even argues, that frequent renewals should take place, because no forward secrecy is provided by this design and an unacknowledged compromise of a private key, this way

Now we will describe, how a user can revoke a compromised private key, e.g. when a device has been lost or hacked.

1. The client, identifying itself with its user credentials, informs the server that the certificate for this username needs to be revoked.
2. The server removes the user's certificate and adds the certificate to the revocation list.
3. If the user still has access to her old private key or the passphrase, she may use it to recover all data from the server and store it locally.
4. The user may then decide to delete the data stored on the server or just remove her own share access and leave the data on the server in case it is shared with another user, that would otherwise lose access.
5. The user generates a new private/public key pair and registers as described in section 6.1. This includes removing the old encrypted private key from the server.
6. The user may decide to re-upload (and re-share) the data she recovered in step 3.

Note, that after a revocation the user loses access to all data shared with her. She might have recovered the data in step 3, if she still had access to her old private key, but she can no longer access the shared folder/file in the cloud. Therefore she will have to ask someone with write access to re-share the data with her (using her new certificate). We consider this necessary, because, if the user herself immediately re-shared all data using her new certificate, an attacker, that had compromised the user's account, might do the same. In that case for other users it would be impossible to know, whether an attacker controls the new certificate or the legitimate user herself. Hence a manual re-share, including a certificate verification, is the more secure approach.

6.4 Sharing data with another user

In this section we will describe, how data is shared with another user. The activity diagram in Figure 3 shows how the algorithm works.

1. Check if the user's X.509 certificate is already known.
 - (a) If the certificate is not known, fetch it from the server. Verify that it is issued by the CA.
 - (b) If the certificate is known, fetch the revocation list from the server. Verify the list is issued by the CA.

would compromise the user's account for all times, allowing a potential attacker to gain access to shares that have long been removed.




Figure 3: Sharing data activity diagram

- i. If the certificate has been revoked, try to fetch a new certificate, verify it is issued by the CA and ask the user to confirm the use of the new certificate.
- ii. If the certificate has not been revoked, continue with 2.
- (c) If no valid certificate is available for the user, inform the user that an encrypted share is currently not possible.
2. Check if the certificate is marked verified.

- (a) If the certificate is marked verified, continue with 3.
 - (b) If the certificate is not marked verified, ask the user to verify it now.
 - i. If the user does not want to verify the certificate now, ask to confirm the use of an unverified certificate.
 - (c) If the user has not confirmed the use of an unverified certificate and no verified certificate is available, inform the user that an encrypted share is currently not possible.
3. Use the certificate to grant access, as described in 5.2.10 and 5.2.11.
- (a) If the used certificate, has been updated (i.e. verified and/or replaced by a new one), update the user's (verified) certificate verification list accordingly.

By using this algorithm, we make sure that the user will 1) only use certificates issued by the CA, 2) be informed if a certificate is revoked and a new one is used,⁴⁶ 3) asked to verify other users' certificates and be informed of the problems with using unverified certificates. At first glance the algorithm may seem complicated and require too much interaction and knowledge from the user, but the great advantage is, that it gives users the choice to decide for themselves, if a certificate change is valid⁴⁷ and if only manually verified certificates should be used or if the user has enough trust in the CA.⁴⁸ Note that usually the sequence is that everything happens automatically except the verification of the certificate, which is why the approach is considered an acceptable trade-off between security and usability. It should be the default, that user's verify each other's certificates before actually sharing data.

7 Public key verification

In our design we chose to base user identity on the possession of a private/public key pair, i.e. a user authenticates herself with her private key. This enables us to use the public key to make cryptographic links and share data only with the owner of the corresponding key. As in most systems that use public key cryptography it is only used for the distribution of symmetric keys (and for digital signatures). While public key cryptography allows us to distribute public keys over potentially insecure channels, which would not be possible with pure symmetric key cryptography, it is prone to man-in-the-middle attacks, where the attacker replaces the public key with her own public key. This means in public key cryptography key verification is a crucial task for the security of the whole system. If we do not get key verification right, it will be easy for an attacker to assume another user's identity and trick others into sharing confidential data with her.

In this section we will discuss the possible solutions to tackle these problems. This involves two aspects, first the key distribution and second the key verification. For

⁴⁶This is important to recognize possible man-in-the-middle attacks.

⁴⁷The user might want to check with her contact.

⁴⁸Optionally, there might be a (non-default) setting, that allows users to turn off the repeated confirmation of the use of unverified certificates.

the process of key distribution we can make leverage of the server as a centralized architecture. This will make it easy to automate key distribution, because no (additional) interaction from the user is required. We will also use a Public Key Infrastructure, in which the server acts as a Certificate Authority (CA), because this makes it hard for any attacker without access to the server to forge a certificate. However, in our threat model we defined the server not to be trustworthy with confidential data and therefore need additional methods for key verification. First we will explain in detail, why the server should not be trusted as a CA. Then we will discuss, Smart Cards, QR Codes and the Socialist Millionaires' Protocol as possible solutions for the problem of key verification.

7.1 Man-in-the-middle attacks

One of the strongest known attacks against public-key crypto systems are man-in-the-middle attacks, which we will describe here, as it is important to understand them for the following section on public key verification. The attacker Mallory is located in the middle between Bob and Alice. Towards Alice Mallory pretends to be Bob and towards Bob he pretends to be Alice, which enables him to intercept all messages exchanged between these two. Such an attack is always possible, when Alice and Bob have no way of verifying who they are actually communicating with. Mallory can just read the messages intercepted and then forward them to their correct receivers, but he might as well not deliver them or first manipulate and then deliver the messages. If the messages carry digital signatures or MACs, however, it is not feasible to manipulate a message from Alice to Bob without Bob noticing, if they have verified each others' keys (Ertel 2012, 90). A man-in-the-middle attack therefore is particularly dangerous, when Alice and Bob have not yet agreed on a secret key or do not have the other one's public key. In this case, when Alice sends her public key to Bob, Mallory can replace it with his own public key, so that Bob will encrypt every message to Alice with Mallory's key. Mallory can then decrypt the messages and, if desired, re-encrypt them with Alice's correct key and forward them to her, so that neither of the two will notice the interception. It is not easy to prevent such an attack, which is why it is so important to verify public keys. In most cases an out-of-band channel (such as a telephone call or a physical meeting between Alice and Bob) is required. Another solution is the use of certificates and a Public Key Infrastructure. All of these solutions have some disadvantages and it depends on the application, which one is best suited to prevent man-in-the-middle attacks (Ferguson 2010, 184f).

7.2 Why the cloud storage provider should not be trusted as a Certificate Authority

As mentioned before we need to consider the possibility that the server acts as a man-in-the-middle and distributes fake keys in order to get access to a user's data. This attack coming from a CSP, that uses a PKI for key verification, where the CSP also acts as the CA has best been described by Wilson and Ateniese (2014, 406f). The man-in-the-middle attack works as follows: When Alice wants to share

a file with Bob, she asks the CSP for Bob’s public key. The CSP then provides a certificate issued to Bob but containing a public key owned by the CSP. Alice will accept that certificate because it has been signed by the CA, i.e. the CSP. When she uses the public key to encrypt the file, the CSP can decrypt it. It may then use Bob’s real public key to re-encrypt it and send it to Bob. So neither Bob nor Alice will ever know of the attack, if the CSP replaced both public keys and forwards all messages after re-encrypting them. This scenario shows how simple, it is for the CSP to spoof any identity it wants, because it authorizes *and* distributes public key certificates used for authentication. In their analysis of the CSPs Wuala, SpiderOak and Tresorit Wilson and Ateniese (2014, 406f) found that all three systems acted as a certificate authority and that these systems were therefore vulnerable to this kind of attack.⁴⁹ However, they could find no evidence any such attack has ever happened. Wilson and Ateniese (2014, 406f) conclude that the CSP cannot be considered a Trusted Third Party and therefore should not act as a CA. Wilson and Ateniese (2014, 411) propose two viable solutions for research: the CSP should 1) allow users to use their own certificates like in a Web of Trust or 2) provide an out-of-band verification mechanism as is the case with off-the-record messaging (OTR). Their findings have been confirmed in a study by Körner et al. (2015, 73), who come to the conclusion that although “the found papers address the issue of data security in general, requirements from the key management have only been expressed rudimentarily.”

From their analysis it becomes clear that we can use the server for key distribution, but cannot solely rely on it for key verification. On the other hand a Web of Trust solution as in PGP requires too much interaction from the users and has too many pitfalls.⁵⁰ We dismiss it, because we want the system to be usable for the average user.

7.3 Smart Card based Public Key Infrastructures

Zwattendorfer et al. (2013) propose to make leverage of the PKI introduced by the Austrian Citizen Card. The core functions of this card are 1) identification and authentication of citizens, 2) generation of qualified digital signatures and 3) encryption of data. For this purpose the citizen card makes use of X.509 certificates and stores the CA’s public key, the user’s certificate and private key on the Smart Card. Any citizen’s key can easily be retrieved from a central public LDAP directory. If Alice in this setting wants to share a file with Bob she will retrieve Bob’s certificate from the LDAP directory, verify the signature with the CA’s certificate stored on her Smart Card and then use Bob’s public key to share data. Note that in order to decrypt the data Bob needs his own Smart Card at hand, because his key is only stored on it and cannot be extracted.

The proposal has some benefits. A Trusted Third Party, the Austrian state,⁵¹ acts

⁴⁹They also mention that Tresorit “does not solely rely on the trustworthiness of the certificate authority issuing the users’ certificates, but also an invitation secret and a pairing password” (Wilson and Ateniese 2014, 411).

⁵⁰Users really need to understand when they should sign a certificate and when they can trust certificates signed by others. Unfortunately this is not a simple task (cf. 3.5.1).

⁵¹This could also be another entity that counts with the required resources to provide similar

as the CA and public keys are accessible to anyone, which makes it easy to integrate it into any cloud storage system. The system can be considered more secure given that encryption is hardware-based. It is much less likely that an attacker manages to extract a user's private key from her Citizen Card than that she compromises a private key stored in software on her device. Because the Smart Card has to be unlocked with a PIN, their system always uses two factor authentication (possession of the card and knowledge of the PIN).

However their design also has some disadvantages. Users always need their Citizen Card and a card reader to access the cloud storage. This sacrifices a lot of usability particularly on mobile devices. More importantly the system transfers trust from the cloud storage provider to the state, because nobody can be certain, that the state has not made a copy of the private key before issuing the Citizen Card. It might also replace certificates in the LDAP, thus performing a similar attack as we described at the beginning of this section. Unfortunately, activists and journalists have good reasons not to put too much trust in the state. Lastly, there are many people on the planet that live in states that have not yet introduced Citizen Cards. This means the system would only be accessible to a small number of people that live in highly developed states.

Because we want to develop software, that does neither put trust in the server nor in an alternative authority, like the state, and also is available to anyone, we dismiss the proposed solution, although it security-wise it is a very interesting approach. In the following sections we will discuss Wilson and Ateniese (2014, 411) second proposal: out-of-band-verification mechanisms.

7.4 QR Codes for Key Exchange and Verification

With the increased use of mobile devices in recent years—a trend that is going to be continued—the question of secure cross-device key exchange mechanisms rises. Any software that is going to be adopted by the users should provide a convenient cross-device key exchange mechanism that does not require too much interaction from the users and yet does not fail to guarantee the authenticity of the secret keys exchanged. However, Körner and Walter (2014, 224) point out, that the use of mobile devices not only brings new challenges but also allows for new forms of interactions we might leverage to facilitate key exchange and key verification. In their paper they propose QR-Codes displayed on one user's devices and scanned by another device as a viable mechanism for key exchange and key verification. Since, in our system it is quite easy to use the server for key exchange, we will focus on key verification.

First note that QR codes count with the necessary capacities to transfer asymmetric keys. The largest QR codes allow to encode 2953 bytes (15624 bits) (Körner and Walter 2014, 225). Considering that the largest keys we might want to exchange are 4096-bit RSA keys, this seems to be enough capacity at first sight. However, it does not suffice to simply transfer the key, but usually we want to transfer a X.509 certificate and check whether it matches the one we obtained from the CSP. This

Smart Cards to all its members.

is also possible with the largest QR codes. Note that QR codes contain checksums that serve to detect errors in the transfer. The basic mechanism is quite simple: If Alice wants to move her private/public key pair to a new device, her client generates a QR code that contains both keys and displays it on the device. The client on the new device scans it, retrieves both keys and after checking the X.509 certificate stores them securely on the new device. If Alice is in an insecure environment she might even encrypt the keys with a password using a function such as PBKDF2. QR codes can also be used for key verification, which is already implemented in mobile messengers like Signal or Threema. If Alice and Bob want to verify each other's key, one of them, say Alice, generates a QR code. This QR code contains the fingerprint of both Alice's and Bob's public key. Bob's client generates the corresponding fingerprints on his device. He scans Alice's QR code and compares the fingerprints generated by Alice and by himself. If they match, Alice and Bob can mark each other's public keys as verified. Of course Alice and Bob can always read out the fingerprints to one another, but it is much more convenient to do it with a simple QR code scan. Still, in case user's are unable to scan fingerprints, the fingerprints can be displayed in characters next to the QR code and used for key verification. While QR codes are a powerful and secure tool on mobile devices that have a camera, they are not very useful for desktop computers and other devices without a camera. It will be impossible to exchange private keys between two desktop computers, e.g. Alice's home computer and her work computer. Therefore we do not consider QR codes a viable solution for cross-device key exchange and instead make leverage of the cloud storage provider for that purpose. However, they can be used to provide key verification. Still, this implies that Alice and Bob have to meet physically and bring their mobile devices in order to verify each other's keys. One disadvantage of QR codes is that they require active physical interaction between the users for key verification before a secure share can take place. Nevertheless we will provide QR codes for key verification in our system. But we are well aware, that we need more mechanism, in case two users cannot physically meet. Also, it remains an important task to encourage users to actually verify each others' keys, because the practice from the messengers shows, that manual key verification is a rare event between average users.

7.5 Socialist Millionaires' Protocol

In this section we will discuss how the Socialist Millionaires' Protocol (SMP) can be used for key verification as an alternative to QR codes. The main benefit of the SMP is, that it does not require the communicating parties to use a (secure) out-of-band channel or meet physically at the time they verify each other's public keys.⁵² However, since it has originally been designed for instant messaging, it requires synchronous communication in order to be secure⁵³ and a direct channel between

⁵²In any key verification it is necessary to exchange information over an out-of-band channel (Körner and Walter 2014, 223f). In SMP this information is represented in the shared secret.

⁵³It would be possible to use the cloud storage provider—adding a layer of indirection—to exchange the messages asynchronously, but that would slow down the process intolerably and thus make it much easier for an attacker to brute force the secret, in which case we would require Alice and Bob to use strong secrets. However, this means sacrificing all the benefits SMP provides regarding

the communicating parties. Both conditions are generally not fulfilled, since all users usually just establish connections with the CSP and not directly to other users.

In a client-side encrypted cloud storage system the SMP needs to be modified as follows to provide key verification. Alice sends a request to the server asking to verify Bob's key. The server passes the request to Bob. If Bob accepts it, before the timeout expires a direct connection between Alice and Bob is established, facilitated by the server. Otherwise Alice will be informed about the timeout and have to retry later. Once the direct connection is established Alice initiates the SMP as described in 3.5.4. When the protocol has terminated successfully, both Alice and Bob mark each other's certificates as verified and tear down the direct connection between them.

7.6 X.509 Certificate and Verification List

In order to keep track of retrieved and verified certificates a "Certificates" folder and "certificate list" file are created in the "Configuration" directory of the client. The folder is used to store the certificates as retrieved from the server. Each entry in the certificate list contains the following information:

- the username
- the verification flag (True/False)
- the fingerprint of the X.509 certificate

The user may decide to synchronize this list across devices by using the usual cloud storage. This means it is encrypted with a unique file key and every write operation to the list has to be signed with a key only the user has access to. When the user logs into her account, she may retrieve the list from the server and update it on her device.⁵⁴

8 Web Access

So far we have discussed how data can be stored, accessed and shared securely via a client software installed on the user's devices. In this section we will discuss, how encrypted data can be accessed via a web panel.

We made it a key requirement for our security, that the client software is open-source, so no back doors can be introduced without the knowledge of the user.⁵⁵ The authenticity and integrity can relatively easy be checked with a digitally signed cryptographic hash of the built. These mechanisms, that make it unlikely that a client software will be modified and malicious code introduced without the knowledge

usability and is therefore dismissed for our application scenario.

⁵⁴In our implementation the synchronization has to be performed manually.

⁵⁵Obviously most users will not really analyze the code, but hopefully the community will. Still, this can never provide absolute security. For example in February 2016 the website of Linux Mint was hacked and the download links were modified to point to a malicious version of the operating system. The integrity checksum on the website were also modified, so that users would not so easily detect the hack. However, the problem was fixed within hours (Linux Mint 2016).

of the user, do not exist for the (JavaScript) applications needed to provide a full functioning web panel.

Given our current design we could easily implement a JavaScript application that fetches the encrypted private key from the server and provided the user's passphrase decrypts the user's files locally. There are two problems with that. First web browsers are generally not a very secure environment and more importantly using JavaScript code delivered by the server to handle the private encryption key means the user has to trust the server not to deliver any malicious code, that will leak information about the data or the encryption keys to the server. For anyone with access to the server such a modification of the code is quite easy, which means a web application actually is a back door by design. Since this problem is well known, Sync (2017b) decided to publish the source code for the web panel while Nextcloud (2017a) decided not to provide a web interface at all. In our threat model we have made the decision not to trust in the server. Therefore we cannot provide full web access to encrypted accounts.

However, interesting approaches were made to find a solution for the problem of JavaScript Code integrity, which is even more relevant when the connection is not encrypted with SSL/TLS. In the later scenario anyone with access to the connection between server and client could modify the JavaScript. What we need is a mechanism for the web browser to check the integrity of the code and if it does not pass not execute it or at least show the user a strong warning message. In 2016 a mechanism to introduce Subresource Identity has been introduced and is supported by the most common web browsers (W3C 2016). Subresource integrity allows developers to specify a cryptographic hash for third party resources in order to protect the integrity. However, when the server itself is not trusted to deliver the correct code, this mechanism is useless, because the browser does not have way to get a cryptographic hash from the server. So there is more research required on how to ensure the integrity of code delivered from the server in security relevant applications.

So far we only know of one solution that actually works: the user has installs an extension in her local browser. This extension has either a fixed integrity checksum for the code or it fetches the integrity checksum from a trusted server, like a Github repository where the complete code is published. The extension will check the code and if it does not pass, display a warning or prevent the execution. A proof of concept of such an extension already exist for the software PrivateBin (2016).

Unfortunately this is far from perfect. The user always has to install an extension each time she wants to access the web interface from a new browser or device. If this is case a user also could simply install the Desktop or mobile application, because the main point of having the web interface is, being able to access one's data from any machine right away without installing any software. It will be hard to keep the extension up to date, if the JavaScript code from the server frequently changes. So the extension might be a hindrance to continuing development and quick bug fixes. Most importantly it cannot be ensured users only access the web interface after they have installed the browser extension. It is not unlikely that users would access the web interface without installing the extension because it is more convenient and the web interface would work without the extension or could at least be modified in such manner.

For these reasons we have decided not to use a web extension to ensure code integrity. It still is not really secure and much less is it usable for the average user. However further research in the field would be welcome.⁵⁶

9 Implementation of the Prototype

In this section we will outline the most important details of our proof of concept implementation. It consists of a server application that is designed to act as a middleware on top of an existing potentially insecure cloud storage, that provides reliability and availability for the stored data. Both the client and server are written in Python 3.5. For cryptography we use the Python module cryptography in both client and server, with the exception of SMP, which is our own implementation, because there are no audited libraries that are suitable for our purposes.

It is important to point out that our applications both client and server are merely a proof of concept. We did not intend to apply secure programming at all points. At many points in our code exceptions and corner cases are not handled appropriately, because we only tested it manually and did not write any unit tests etc.. Therefore our application should not be used in a production environment. Another reason is, that both applications currently do not store private keys securely in the Operating System's key ring (or by setting appropriate file access rights) but simply store keys as a file in the application folder. For more details about these limitations have a look at our TU-Berlin Gitlab Repository, where we also published our code.⁵⁷

9.1 Server

The main purpose of the server application is on the one hand to act as a CA and the key distributor in a PKI and on the other hand store the users' confidential data. The server is a small Flask application of about 1000 lines of code. The server provides a RESTful API which the client uses to retrieve and modify information. The server has access to a Mysql database, which stores all information on user accounts, certificates and revocations for the PKI and even the user data itself. The database and especially user account data is also used to manage access control.

We do not use a real cloud storage in our proof of concept, but only store the users' data in a database.⁵⁸ This is acceptable, because our focus is mostly on the client, where the encryption/ decryption operations happen, and it is not at all on the proving the usage and security guarantees that cloud storages provide. The server also stores notifications, which clients can retrieve in order to be informed of new shared data with the user or new SMP certificate verification requests.

⁵⁶As a security-usability trade-off one might at least want to allow users to share certain folders via web link. This is a relatively simple task given the tree structure, we introduced. Tresorit (2017b) has designed a very convenient solution for encrypted web links. However, we do not want to include it in our design for the concerns mentioned in this section.

⁵⁷https://gitlab.tubit.tu-berlin.de/gaspar_ilom/crypt-cloud

⁵⁸For this proof of concept implementation we can store files up to 4GB of size (after encryption).

9.2 Client

The client application consists of about 1700 lines of code. Currently we only implemented a Desktop client, but there is no reason not to port our code to provide mobile clients. Our application uses Python's requests module for connection's with the server. After login (and registration) is offers a main menu, in which the user can choose between several options regarding file management and certificate verification. It also constantly requests notifications from the server in order to inform the user of new data accessible or new verification requests.

The client does not implement the complete cryptographic read and write trees as proposed in 5. Since similar cloud storages are already working in practice in our proof of concept we focus on the key verification and key exchange mechanisms instead. That means sharing and uploading files to the cloud in our implementation is only a mockup we implemented in order to prove that our PKI, key exchange and key verification mechanisms really provide the cryptographic means to share data securely. The user is only able to upload individual files and has to do so manually. Before uploading files the file data and their names are encrypted with AES-128-CBC using a PKCS#7 padding and HMAC-SHA256 for authentication. Because both AES and HMAC are based on the same key, there is no separation of write and read access in this proof of concept. The 128 bit key itself is encrypted with the user's private key and also uploaded to the server. For sharing it can be re-encrypted with another user's private key and attached to the original key data on the server. This is the equivalent of what we earlier referred to as a cryptographic link. In our implementation it is possible to revoke a share by removing the encrypted key for a user from the server,⁵⁹ but we did not implement lazy re-encryption as proposed in section 5, so the actual key used for encryption will not be changed.⁶⁰ Revocation of a share does not require access to a user's private key or a revocation certificate, but only to the the user's login credentials (username/email and password). In production this should be changed, because otherwise an attacker might revoke a user's certificate, in order to forge a new one controlled by her.⁶¹

When adding the initial device, the client registers an account and then generates a private key as explained in section 6.1. The private key is encrypted with the random generated 12 words passphrase, according to the PKCS#8 standard. This means PBKDF2 with 2048 iterations of HMAC-SHA256 and a 64 bit random salt is used to derive a key and use AES-256-CBC to encrypt the private key with the derived key. We point out that we use cryptography's BestAvailableEncryption class, so that this algorithm might change over time. After encrypting it the private key is stored both on the server and the user's device. The clients also requests a corresponding X.509 certificate from the server. The twelve word passphrase is displayed to the user (to remember it or write it down and store it in a safe place). It is also stored

⁵⁹The server also implements an ACL, in order to control who has access to a file.

⁶⁰The only way to change the key currently is to re-upload the file (and re-share it, if necessary).

⁶¹On the other hand requiring a signature from the private key or a revocation certificate means that some users will not be able to revoke their private keys, which also represents a security threat. So in this case it should really be ensured that users store their revocation certificates in a safe place, they will not lose access to.

on the device.⁶² If the user ever loses her device or wants to add a new one, she can retrieve the key from the server and decrypt it with the twelve word passphrase as explained in section 6.2.

Once the the client is configured, i.e. the user registered, logged in and has a valid certificate and corresponding private key, she can upload data to the cloud and share it. For sharing we implemented our algorithm as explained in section 6.4. In this algorithm the user is asked to verify the counterparts certificate. For doing so, we provide two options: QR-Codes and the Socialist Millionaires' Protocol.

The activity diagram in Figure 4 shows how QR-Code verification works. The client uses the fingerprints of both users' certificates in order to compute a hash value. If the user wants to display the QR-Code this value is Base64 encoded before generating and displaying the QR-Code. The Base64 encoding is necessary, because the Python library used to scan the QR-Codes `zbar-py` has some issues when scanning non-ASCII characters. Since the SHA256 hash value is only 32 bytes long this is not a problem. After displaying the QR-Code the user is asked, whether the scanning on her counterpart's device succeeded. If so, she can answer yes and the counterparts certificate is marked as verified. This is possible because we use both users' certificates to compute the fingerprint and therefore the verification is always mutual.⁶³ If the user scans the QR-Code, the client will read images from the camera of the device and try to find a QR-Code. Once it has found a QR-Code, it decodes the value read and compares it to the hash value computed earlier. If the values match, the user can be certain, that the same certificates are used on both devices.

In order to implement the Socialist Millionaires' Protocol we modified Shane Tully's (2013) version according to our needs. Most importantly we converted his code to Python3 and changed the secret generation. In our code the secret is a SHA256-hash computed over a concatenation of the SHA256-fingerprints of both participants' certificates and the shared secret. We point out that the shared secret and thus the computed secret are never transmitted over a network. Our clients do not open direct connections to one another, but relay there messages to the server. The complete interaction, as we implement it, is displayed in Figure 5. Each of the messages posted to the API provided by the server is Base64 encoded, in order to be more easily handled by the flask-RESTful module.⁶⁴

We considered opening a peer to peer connection between the clients in order to perform the SMP verification, but finally decided against that. It is not necessary, because the SMP was created exactly for a scenario in which the server might act as a man in the middle.⁶⁵ It ensures, that such behavior will be detected. The SMP

⁶²In production the operating system's key ring should be used for that purpose.

⁶³Theoretically it is possible that the counterpart uses a manipulated version of the client application, that will always display a successful result. That way the user could be tricked into verifying the counterpart's certificate. However, if the counterpart's client really was manipulated by an attacker, the attacker would probably have access to all files and the user's private key anyway. So this does not represent a weakness an attacker could exploit to gain access to new data.

⁶⁴The computational and spatial overhead is not considered relevant for a proof of concept application, but the Base64 encoding could be avoided if required to enhance performance (see section 10).

⁶⁵It is used in a similar context in OTR, where an XMPP server might act as a man in the




Figure 4: QR-Code verification activity diagram

Note that *value* actually is calculated from the SHA256 fingerprints of both certificates.
QR-Codes are Base64 encoded.

requires a synchronous session and we arbitrarily set a time limit to one minute. If the protocol has not terminated one minute after the client posted the first cryptographically relevant data, it is aborted and the verification fails. Further research is required in order to figure out, what would be a time limit that on the one hand makes sure a man in the middle does not have enough time to brute force a weak natural language secret and on the other hand still provide enough time for the client to successfully finish the protocol. It is important to point out that neither our implementation nor Shane Tully's have been security audited. So it should not be used for critical purposes.

middle.

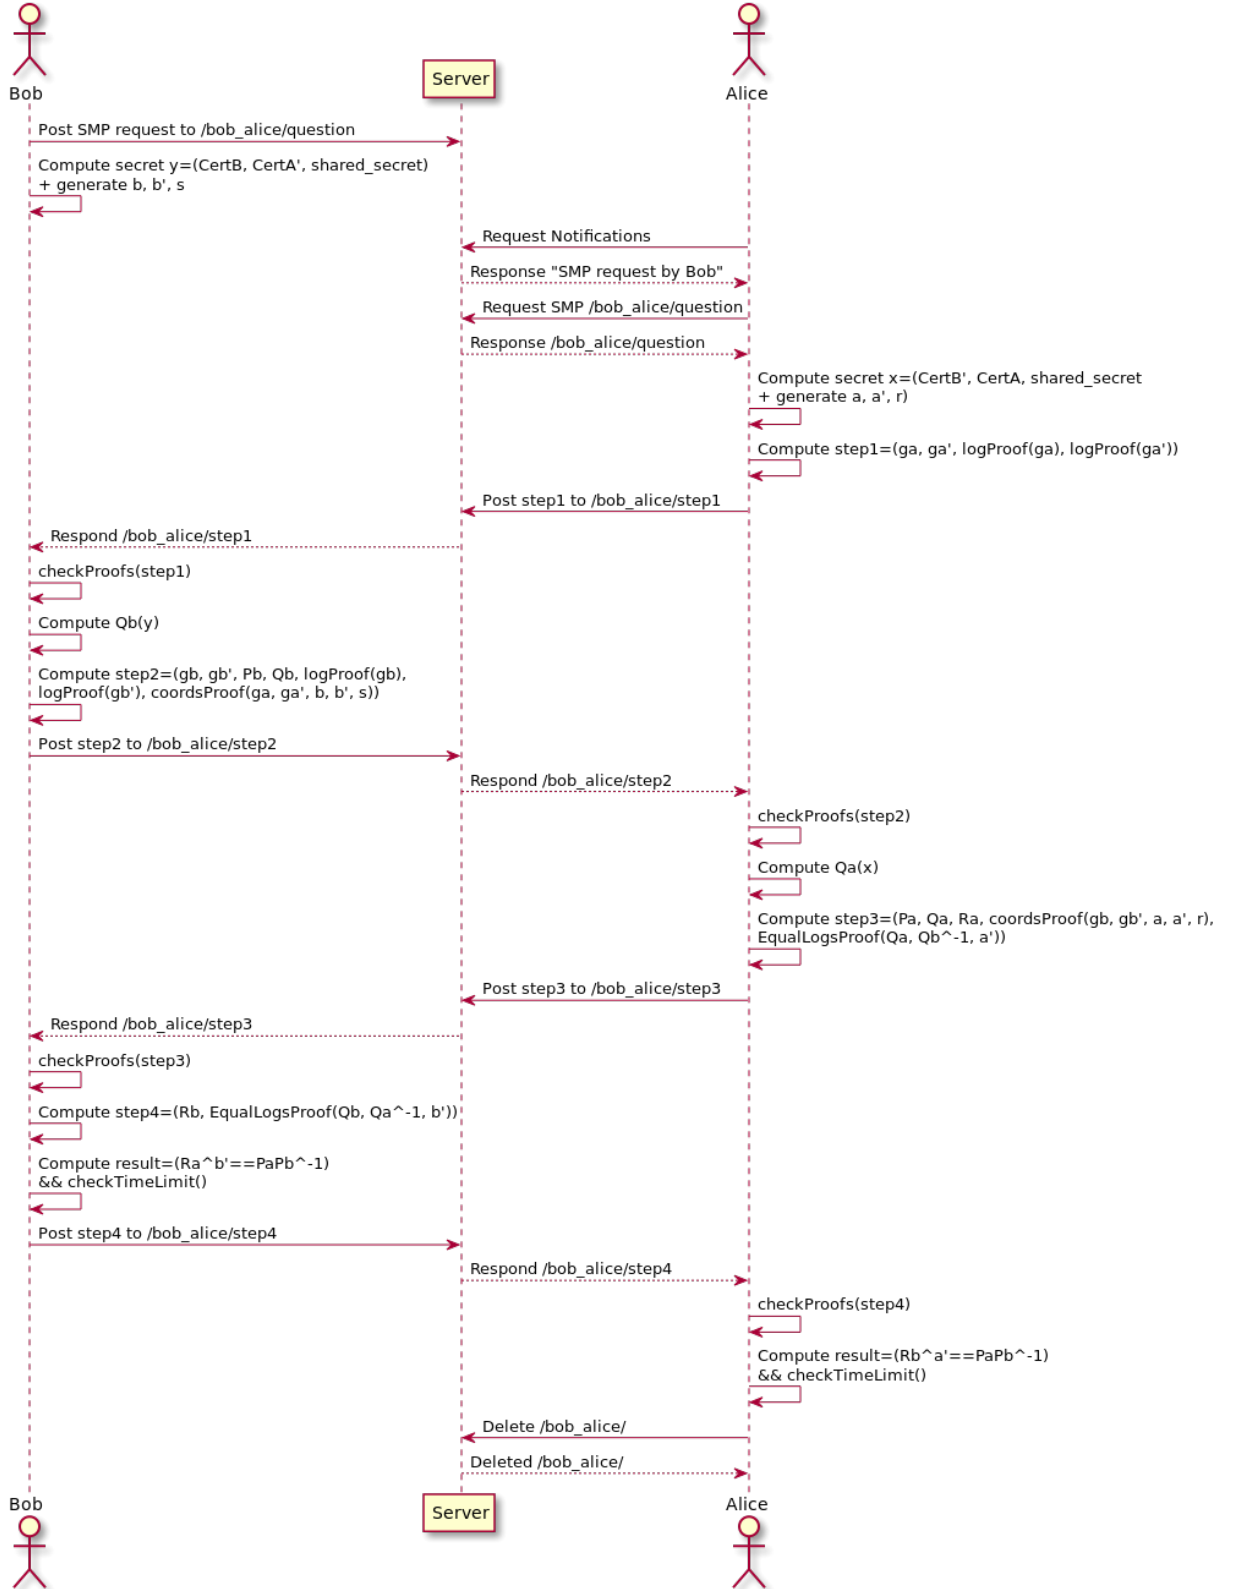


Figure 5: SMP verification sequence diagram

Note that not all requests by the clients to the server are included, because the diagram is already complex enough.

10 Evaluation of the Prototype

In this section we will briefly evaluate the prototype. With our proof of concept we showed, that the system we designed really works and allows users to share end-to-end encrypted data over an untrusted server. However, there are some limitations. Our goal was to provide a cloud storage system, that allows users to share data with client-side encryption. We achieved that with our proof of concept implementation. We even proposed a cryptographic key structure, that is more efficient and powerful than what we implemented. Other evaluations, e.g. by Grolimund et al. (2006) have shown the efficiency of similar cryptographic tree structures. Due to the cryptographic tree structure the use of computing power and bandwidth is limited to a minimum. For instance granting access does not require any re-encryption operations. Only one cryptographic link is added to the tree, because the user inherits all downward access rights from this point.

We implemented a mechanism that allows users to exchange their private key between devices by storing it encrypted on the server and retrieve it from there. All they need is to remember a 12 English words passphrase. Alternatively they can display it on one of the devices, they already added to the account. 12 randomly chosen words provide enough entropy not to be easily brute forced by an attacker. However, it might also be too long to be remembered by the average user. Further research is needed regarding the usability of this key exchange mechanism, i.e. whether users actually find this acceptable or whether better solutions to enforce security are needed. We can state that, given the security of our encryption schemes, we found a secure way to transfer confidential key material between a users devices without requiring much interaction from the user.

We also developed a small PKI, in which the server application acts as CA and the key distributor. Since PKIs are well-proven concepts, we will not further evaluate it. Once a user has retrieved another user's certificate from the PKI, she can verify it either using QR-Code verification or SMP verification. Our proof of concept shows that both models do not only work for instant messaging, but are also suitable for the context of cloud storages. They provide technically secure mutual certificate verification. Yet QR-Codes usually require users to meet physically. Although QR-Codes might also be shared over an out-of-band channel, so that no physical meeting is necessary. The QR-Code verification is actually simply an automating of the old fashioned fingerprint comparison.

The Socialist Millionaires' Protocol in contrast neither requires a physical meeting nor an out-of-band channel for certificate verification, but it requires a synchronous session in order to be secure. Otherwise we could not use a weak natural language secret, because an attacker might have enough time to brute force the secret. Our implementation works reliably. Yet again we have to ask the question, whether users actually trust the protocol (which is too complicated to be understood by most of them) and find it acceptable to require a simultaneous session to verify their certificates. SMP and QR-Code verification well complement each other according to the users' needs and preferences. Even if users do not verify other users' certificates they are still able to share data with them. Because this means they fully trust the CSP as a CA, we implemented our sharing logic such, that it strongly encourages

users to actually perform a certificate verification. Again this might be deterring for some users, but hopefully not for high profile users (like activists and journalists) that are aware of the threat an insecure cloud storage represents.

We have proven that our concepts technically work, now it depends on, whether users actually like them. Let us finish with a few words regarding the performance of our system and, why we decided not to run benchmarks. We believe that performance is not an issue for the system we designed for the following reasons. The innovative parts of our system mostly concern the mechanism to exchange keys between devices and verification of other users' certificates. QR-Code verification does not require any interaction with the server at all and only two users will verify each others keys at a time by transmitting data exclusively within the QR-Code. No data is stored permanently on either device. In regard to SMP verification the server has to store a natural language question and four messages which range from 428 to 1676 bytes. In our implementation the messages are Base64 encoded before sending, which means 3 bytes will be 4 bytes effectively in transmission and storage at the server. So the maximum message size is 2235 bytes to be transmitted and stored on the server. Considering that the server should be able to retrieve and deliver large files of hundreds of MBs and that the SMP data will only be stored for less than a minute, there is no problem regarding performance. Note, that the server does not have to perform any calculations. The clients on the other hand have to calculate two of these messages per SMP verification, but computationally this is more or less equivalent to several Diffie-Hellman Key Exchanges (because of the modular exponentiations involved) so this also scales.

The user's private keys and certificates (one of each per user) are just a few thousand bytes large and are not an issue to be stored on the database. They will only be accessed when registering a new account or adding a new device to that account. So it should happen much less frequently than the usual data access. Summarizing on the performance of the system we conclude that our key exchange and verification mechanisms will not be the bottleneck of the cloud storage neither at client- nor at server-side.

11 Conclusion

This thesis presented a new client-side encrypted cloud storage system and successfully proved its feasibility in the evaluation of a proof of concept. We introduced an improved version of the cryptographic tree structure Cryptree, which enables us to flexibly and intuitively share data and distinguish between read and write access. An encryption key is always just used for a single file, which reduces the amount of necessary encryption/decryption operations to a minimum, and thus provides the basis for an efficient client-side encrypted cloud storage. Future research could go even further and find ways of using functional cryptography in order to avoid the re-encryption and re-upload of a complete file after each modification while also providing file versioning and simultaneous modification of encrypted documents by multiple users.

We presented a cross-device key exchange mechanism that allows users to flexibly add (and remove) devices to their account and even recover from the loss of all

devices. The innovation is that this is possible without a security trade-off. Because the user's private key is encrypted with a random generated 12 words passphrase and stored in the cloud, it is ensured that the encryption key used to protect the private key has enough entropy and still is easily recoverable for the user. This does not only protect the confidentiality of the data in the cloud, but also its availability, because users will not easily lose access to their devices.

In order to make sharing easy to use a PKI, run by the CSP, provides access to all users' certificates. Apart from sharing this also enables easy to use certificate revocation and the removal of particular shares. However, contrary to existing solution our design does not require any trust in the server, because it provides efficient and easy to use means of certificate verification. This is achieved by either QR-Code scanning or the Socialist Millionaires' Protocol. While the former requires users to physically meet or send the QR-Code over an out-of-band channel to mutually verify their certificates, the latter can be performed over an insecure channel. This insecure channel may even be controlled by the CSP itself, as it is the case in our implementation. All that is required from the users is a synchronous session and a shared secret, that can even be a dictionary word.

The client-side encrypted cloud storage system we designed is very usable. It preserves the complete Sync & Share functionality with the exception of web access, which is not included, because there is no way to provide without implicitly requiring users to trust the code for decryption in the browser, which is delivered by the untrusted CSP. Our design shall be published under an open source license, so that anyone willing is able to verify the client application really works as we claim. Due to the newly introduced features we were able to present a new client-side encrypted cloud storage system that exceeds the degree of security of the users' data guaranteed by former solutions. Therefore our system is particularly suitable for users, that must heavily rely on the confidentiality, integrity and authenticity of their data such as activists, lawyers, journalists and others.

References

- Amazon. 2018. *Amazon S3 Reduced Redundancy Storage*. Visited on 03/16/2018. <https://aws.amazon.com/s3/reduced-redundancy/>.
- Bourrier, Mathieu. 2016. *ownCrypt_CryptoDesign-1*. Visited on 06/23/2017. <https://cloud.airmail.fr/s/d475f3c9d77e99d8b413977729be5f0e/download>.
- . 2015. *ownCrypt_Intents&ThreatModel&Features-2*. Visited on 06/23/2017. <https://cloud.airmail.fr/s/a9fe1a55a59374dd6bc399f4ff4d60db/download>.
- Cachin, Christian, Idit Keidar, and Alexander Shraer. 2009. “Trusting the Cloud”. *SIGACT News* 40 (2): 81–86.
- Cryptomator. 2017a. *Cryptomator - Homepage*. Visited on 09/21/2017. <https://cryptomator.org>.
- . 2017b. *Cryptomator - Sicherheitsarchitektur*. Visited on 09/21/2017. <https://cryptomator.org/architecture/>.
- Delfs, Hans, and Helmut Knebl. 2015. *Introduction to Cryptography*. Information Security and Cryptography. DOI: 10.1007/978-3-662-47974-2. Berlin, Heidelberg: Springer Berlin Heidelberg.
- Diffie, W. 1992. “The First Ten Years of Public-Key Cryptology”. In *Contemporary Cryptology: The Science of Information Integrity*, 135–175. IEEE Press New York.
- Diffie, W., and M. Hellman. 1976. “New directions in cryptography”. *IEEE Transactions on Information Theory* 22 (6): 644–654.
- Ertel, Wolfgang. 2012. *Angewandte Kryptographie*. München: Hanser.
- Ferguson, Niels. 2010. *Cryptography engineering : design principles and practical applications*. Indianapolis, Ind.: Wiley.
- Frosch, T., C. Mainka, C. Bader, F. Bergsma, J. Schwenk, and T. Holz. 2016. “How Secure is TextSecure?” In *2016 IEEE European Symposium on Security and Privacy (EuroS P)*, 457–472.
- Gadhavi, L., M. Bhavsar, M. Bhatnagar, and S. Vasoya. 2016. “Design of efficient algorithm for secured key exchange over Cloud Computing”. In *2016 6th International Conference - Cloud System and Big Data Engineering (Confluence)*, 180–187.
- Grolimund, D., L. Meisser, S. Schmid, and R. Wattenhofer. 2006. “Cryptree: A Folder Tree Structure for Cryptographic File Systems”. In *2006 25th IEEE Symposium on Reliable Distributed Systems (SRDS’06)*, 189–198.
- Hallberg, Sven Moritz. 2008. “Individuelle Schlüsselverifikation via Socialist Millionaires’ Protocol”. Visited on 09/23/2017. <http://www.khjk.org/log/2012/apr/smp/smp.pdf>.
- IETF. 2008. *RFC 5280: Internet X.509 Public Key Infrastructure Certificate and Certificate Revocation List (CRL) Profile*. Visited on 10/11/2017. <https://tools.ietf.org/html/rfc5280>.

- Kamara, Seny, and Kristin Lauter. 2010. "Cryptographic Cloud Storage". In *Financial Cryptography and Data Security*, 136–149. Springer, Heidelberg.
- Kaur, Manpreet, and Rajbir Singh. 2013. "Implementing encryption algorithms to enhance data security of cloud in cloud computing". *International Journal of Computer Applications* 70 (18).
- Kim, Tiffany Hyun-Jin, Lin-Shung Huang, Adrian Perrig, Collin Jackson, and Virgil Gligor. 2013. "Accountable Key Infrastructure (AKI): A Proposal for a Public-key Validation Infrastructure". In *Proceedings of the 22Nd International Conference on World Wide Web*, 679–690. WWW '13. Rio de Janeiro, Brazil: ACM.
- Kiss, Jemima. 2014. "Snowden: Dropbox is hostile to privacy, unlike 'zero knowledge' Spideroak". *The Guardian* (London). Visited on 09/21/2017. <http://www.theguardian.com/technology/2014/jul/17/edward-snowden-dropbox-privacy-spideroak>.
- Körner, Kevin, Holger Kühner, Julia Neudecker, Hannes Hartenstein, and Thomas Walter. 2015. "Bewertungskriterien und ihre Anwendung zur Evaluation und Entwicklung sicherer Sync&Share-Dienste." In *DFN-Forum Kommunikationstechnologien*, 71–82. Visited on 09/27/2017. <http://cs.emis.de/LNI/Proceedings/Proceedings243/71.pdf>.
- Körner, Kevin, and Thomas Walter. 2014. "Sichere und benutzerfreundliche Schlüsselverteilung auf Basis von QR-Codes." In *GI-Jahrestagung*, 223–234.
- Kumar, A., B. G. Lee, H. Lee, and A. Kumari. 2012. "Secure storage and access of data in cloud computing". In *2012 International Conference on ICT Convergence (ICTC)*, 336–339.
- Lám, I., S. Szebeni, and L. Buttyán. 2012a. "Invitation-Oriented TGDH: Key Management for Dynamic Groups in an Asynchronous Communication Model". In *2012 41st International Conference on Parallel Processing Workshops*, 269–276.
- . 2012b. "Tresorium: Cryptographic File System for Dynamic Groups over Untrusted Cloud Storage". In *2012 41st International Conference on Parallel Processing Workshops*, 296–303.
- Lám, István. 2016. *Was ist Zero-Knowledge-Verschlüsselung?* Visited on 09/17/2017. <https://tresorit.com/blog/zero-knowledge-verschlüsselung/>.
- Law, Effie Lai-Chong, Virpi Roto, Marc Hassenzahl, Arnold P.O.S. Vermeeren, and Joke Kort. 2009. "Understanding, Scoping and Defining User Experience: A Survey Approach". In *Proceedings of the SIGCHI Conference on Human Factors in Computing Systems*, 719–728. CHI '09. Boston, MA, USA: ACM.
- Linux Mint. 2016. *Beware of hacked ISOs if you downloaded Linux Mint on February 20th! – The Linux Mint Blog*. Visited on 10/16/2017. <https://blog.linuxmint.com/?p=2994>.

- Meister, Andre. 2017. "Staatstrojaner: Bundestag hat das krasseste Überwachungsgesetz der Legislaturperiode beschlossen (Updates)". *netzpolitik.org*. Visited on 10/21/2017. <https://netzpolitik.org/2017/staatstrojaner-bundestag-beschliesst-diese-woche-das-krasseste-ueberwachungsgesetz-der-legislaturperiode/>.
- Melara, Marcela S., Aaron Blankstein, Joseph Bonneau, Edward W. Felten, and Michael J. Freedman. 2015. "CONIKS: Bringing Key Transparency to End Users". In *24th USENIX Security Symposium (USENIX Security 15)*, 383–398. Washington, D.C.: USENIX Association.
- Nextcloud. 2017a. *End-to-end encryption design*. Visited on 10/01/2017. <https://nextcloud.com/endtoend/>.
- . 2017b. *Nextcloud*. Visited on 10/01/2017. <https://nextcloud.com/>.
- Nielsen, Jakob. 2017. *Usability 101: Introduction to Usability*. <https://www.nngroup.com/articles/usability-101-introduction-to-usability/>.
- OTR Development Team. 2017. *Off-the-Record Messaging Protocol version 3*. Visited on 10/12/2017. <https://otr.cypherpunks.ca/Protocol-v3-4.0.0.html>.
- Pancholi, Vishal R, and Bhadresh P Patel. 2016. "Enhancement of cloud computing security with secure data storage using AES". *International Journal for Innovative Research in Science and Technology* 2 (9): 18–21.
- Percival, Colin. 2009. "Stronger key derivation via sequential memory-hard functions", 1–16. BSD Can. Visited on 09/27/2017. http://www.bsdcan.org/2009/schedule/attachments/87_scrypt.pdf.
- PrivateBin. 2016. *WebExtension: A proof-of-concept webextension for verifying JS integrity*. Original-date: 2016-12-11T18:36:50Z. Visited on 10/16/2017. <https://github.com/PrivateBin/WebExtension>.
- Raymond. 2015. *SafeMonk: Dropbox + Security Made Simple*. Visited on 09/22/2017. <http://www.dropboxwiki.com/dropbox-addons/safemonk-dropbox-security-made-simple>.
- Rewagad, P., and Y. Pawar. 2013. "Use of Digital Signature with Diffie Hellman Key Exchange and AES Encryption Algorithm to Enhance Data Security in Cloud Computing". In *2013 International Conference on Communication Systems and Network Technologies*, 437–439.
- Rivest, R. L., A. Shamir, and L. Adleman. 1978. "A method for obtaining digital signatures and public-key cryptosystems". *Communications of the ACM* 21 (2): 120–126. Visited on 09/15/2017. <http://portal.acm.org/citation.cfm?doid=359340.359342>.
- Rueppel, Rainer A. 1992. "Stream Ciphers". In *Contemporary Cryptology: The Science of Information Integrity*, ed. by G.J. Simmons, 65–134. IEEE Press New York.
- SafeMonk. 2013. *Key Management and Organization*. Visited on 08/13/2013. <http://support.safemonk.com/customer/portal/articles/1018901-key-management-and-organization>.

- Scherschel, Fabian A. 2015. "Adieu sichere Dropbox: Cloud-Verschlüsseler SafeMonk gibt auf". *heise Security*. Visited on 09/22/2017. <http://www.heise.de/security/meldung/Adieu-sichere-Dropbox-Cloud-Verschluesseler-SafeMonk-gibt-auf-2637330.html>.
- Schneier, Bruce. 1996. *Applied Cryptography : Protocols, Algorithms, and Source Code in C*. 2. ed. New York: Wiley.
- Seafle. 2017. *Security features · Seafle Server Manual*. Visited on 10/01/2017. https://manual.seafile.com/security/security_features.html.
- Sookasa. 2017. *SOOKASA WHITEPAPER / SECURITY*. Visited on 09/22/2017. https://www.sookasa.com/wp-content/uploads/2016/01/Jan2016_Sookasa_WhitePaper_Security.pdf.
- SpiderOak. 2017a. *Encryption White Paper*. Visited on 09/21/2017. <https://spideroak.com/resources/encryption-white-paper>.
- . 2017b. *No Knowledge, Secure-by-Default Products*. Visited on 09/17/2017. <https://spideroak.com/no-knowledge/>.
 - . 2017c. *Why We Will No Longer Use the Phrase Zero Knowledge to Describe Our Software*. Visited on 09/17/2017. <https://spideroak.com/articles/why-we-will-no-longer-use-the-phrase-zero-knowledge-to-describe-our-software/>.
- Swoboda, Joachim, Stephan Spitz, and Michael Pramateftakis. 2008. *Kryptographie und IT-Sicherheit : Grundlagen und Anwendungen*. 1. Aufl. Studium : IT-Sicherheit und Datenschutz. Wiesbaden: Vieweg + Teubner.
- Sync. 2017a. *Privacy Matters*. Visited on 09/26/2017. <https://www.sync.com/your-privacy/>.
- . 2017b. *Privacy White Paper*. Visited on 09/26/2017. <https://www.sync.com/pdf/sync-privacy.pdf>.
- Syncplicity. 2015. *Syncplicity Security and Control Features*. Visited on 09/21/2017. <https://www.syncplicity.com/resources/datasheets/syncplicity-security-and-control-features>.
- T., Finanzmagazin I. 2014. *Zero Knowledge wird zum Datentransfer-Gütesiegel · IT Finanzmagazin*. Visited on 09/21/2017. <https://www.it-finanzmagazin.de/zero-knowledge-wird-zum-datentransfer-guetesiegel-1888/>.
- Tresorit. 2017a. *Tresorit - Security*. Visited on 09/26/2017. <https://tresorit.com/security>.
- . 2017b. *White Paper*. Visited on 09/26/2017. <https://tresorit.com/files/tresoritwhitepaper.pdf>.
- Tully, Shane. 2013. *MITM Protection via the Socialist Millionaire Protocol (OTR-style)*. Blog. Visited on 03/13/2018. <https://shanetully.com/2013/08/mitm-protection-via-the-socialist-millionaire-protocol-otr-style/>.
- W3C. 2016. *Subresource Integrity*. Visited on 10/16/2017. <https://www.w3.org/TR/SRI/>.

- Whitten, Alma, and J. D. Tygar. 1998. *Usability of Security: A Case Study*, tech. rep. CMU-CS-98-155. CARNEGIE-MELLON UNIV PITTSBURGH PA DEPT OF COMPUTER SCIENCE, CARNEGIE-MELLON UNIV PITTSBURGH PA DEPT OF COMPUTER SCIENCE. Visited on 09/29/2017. <http://www.dtic.mil/docs/citations/ADA361032>.
- Wilson, Duane C., and Giuseppe Ateniese. 2014. "To Share or not to Share" in Client-Side Encrypted Clouds". In *Information Security: 17th International Conference, ISC 2014, Hong Kong, China, October 12-14, 2014. Proceedings*, ed. by Sherman S. M. Chow, Jan Camenisch, Lucas C. K. Hui, and Siu Ming Yiu, 401–412. Cham: Springer International Publishing.
- Yao, F.F., and Y.L. Yin. 2005. "Design and Analysis of Password-Based Key Derivation Functions". *IEEE Transactions on Information Theory* 51 (9): 3292–3297. Visited on 09/14/2017. <http://ieeexplore.ieee.org/document/1499059/>.
- Zhao, G., C. Rong, J. Li, F. Zhang, and Y. Tang. 2010. "Trusted Data Sharing over Untrusted Cloud Storage Providers". In *2010 IEEE Second International Conference on Cloud Computing Technology and Science*, 97–103.
- Zwattendorfer, Bernd, Bojan Suzic, Peter Teufl, and Andreas Derler. 2013. "Sicheres Speichern in der Public Cloud mittels Smart Cards". In *D-A-CH Security*, 120–132. Nürnberg.